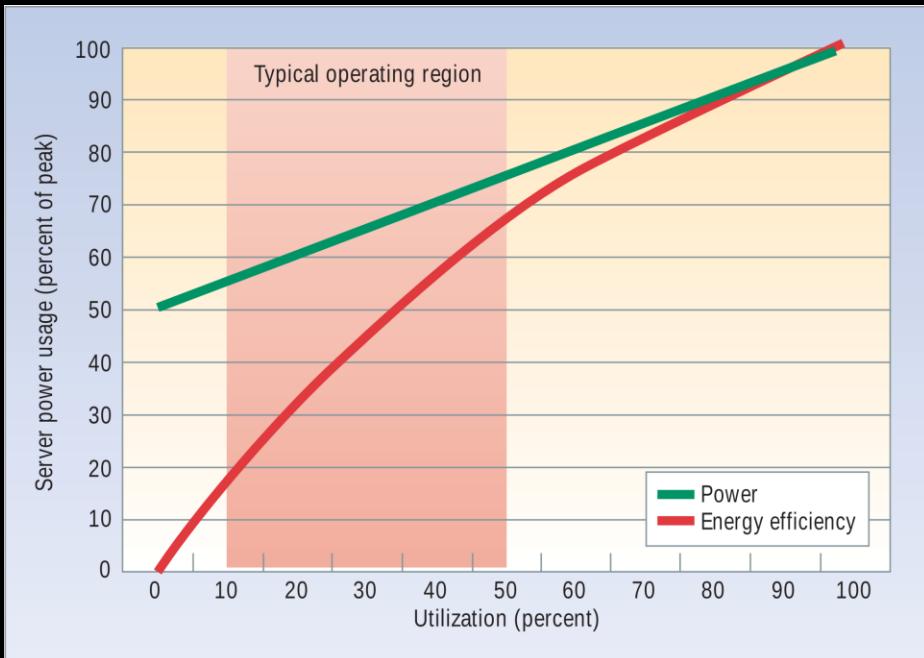


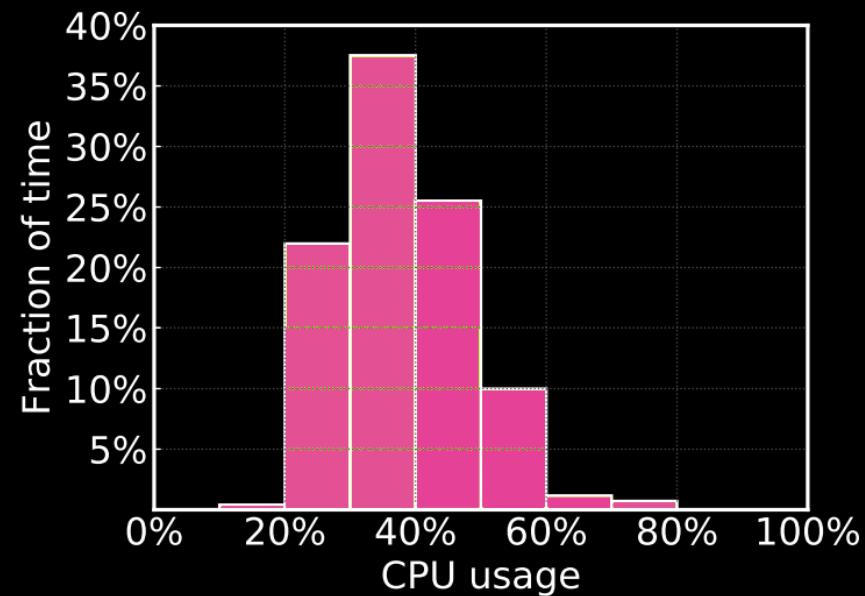
Give our Lambda's Connectivity, or else...

...we'll take it for ourselves

Understanding the Problems Serverless wants to Solve



Machine **under-utilization** is **energy-inefficient**



Machines in datacenters are **very underutilized**

Understanding the Problems Serverless wants to Solve

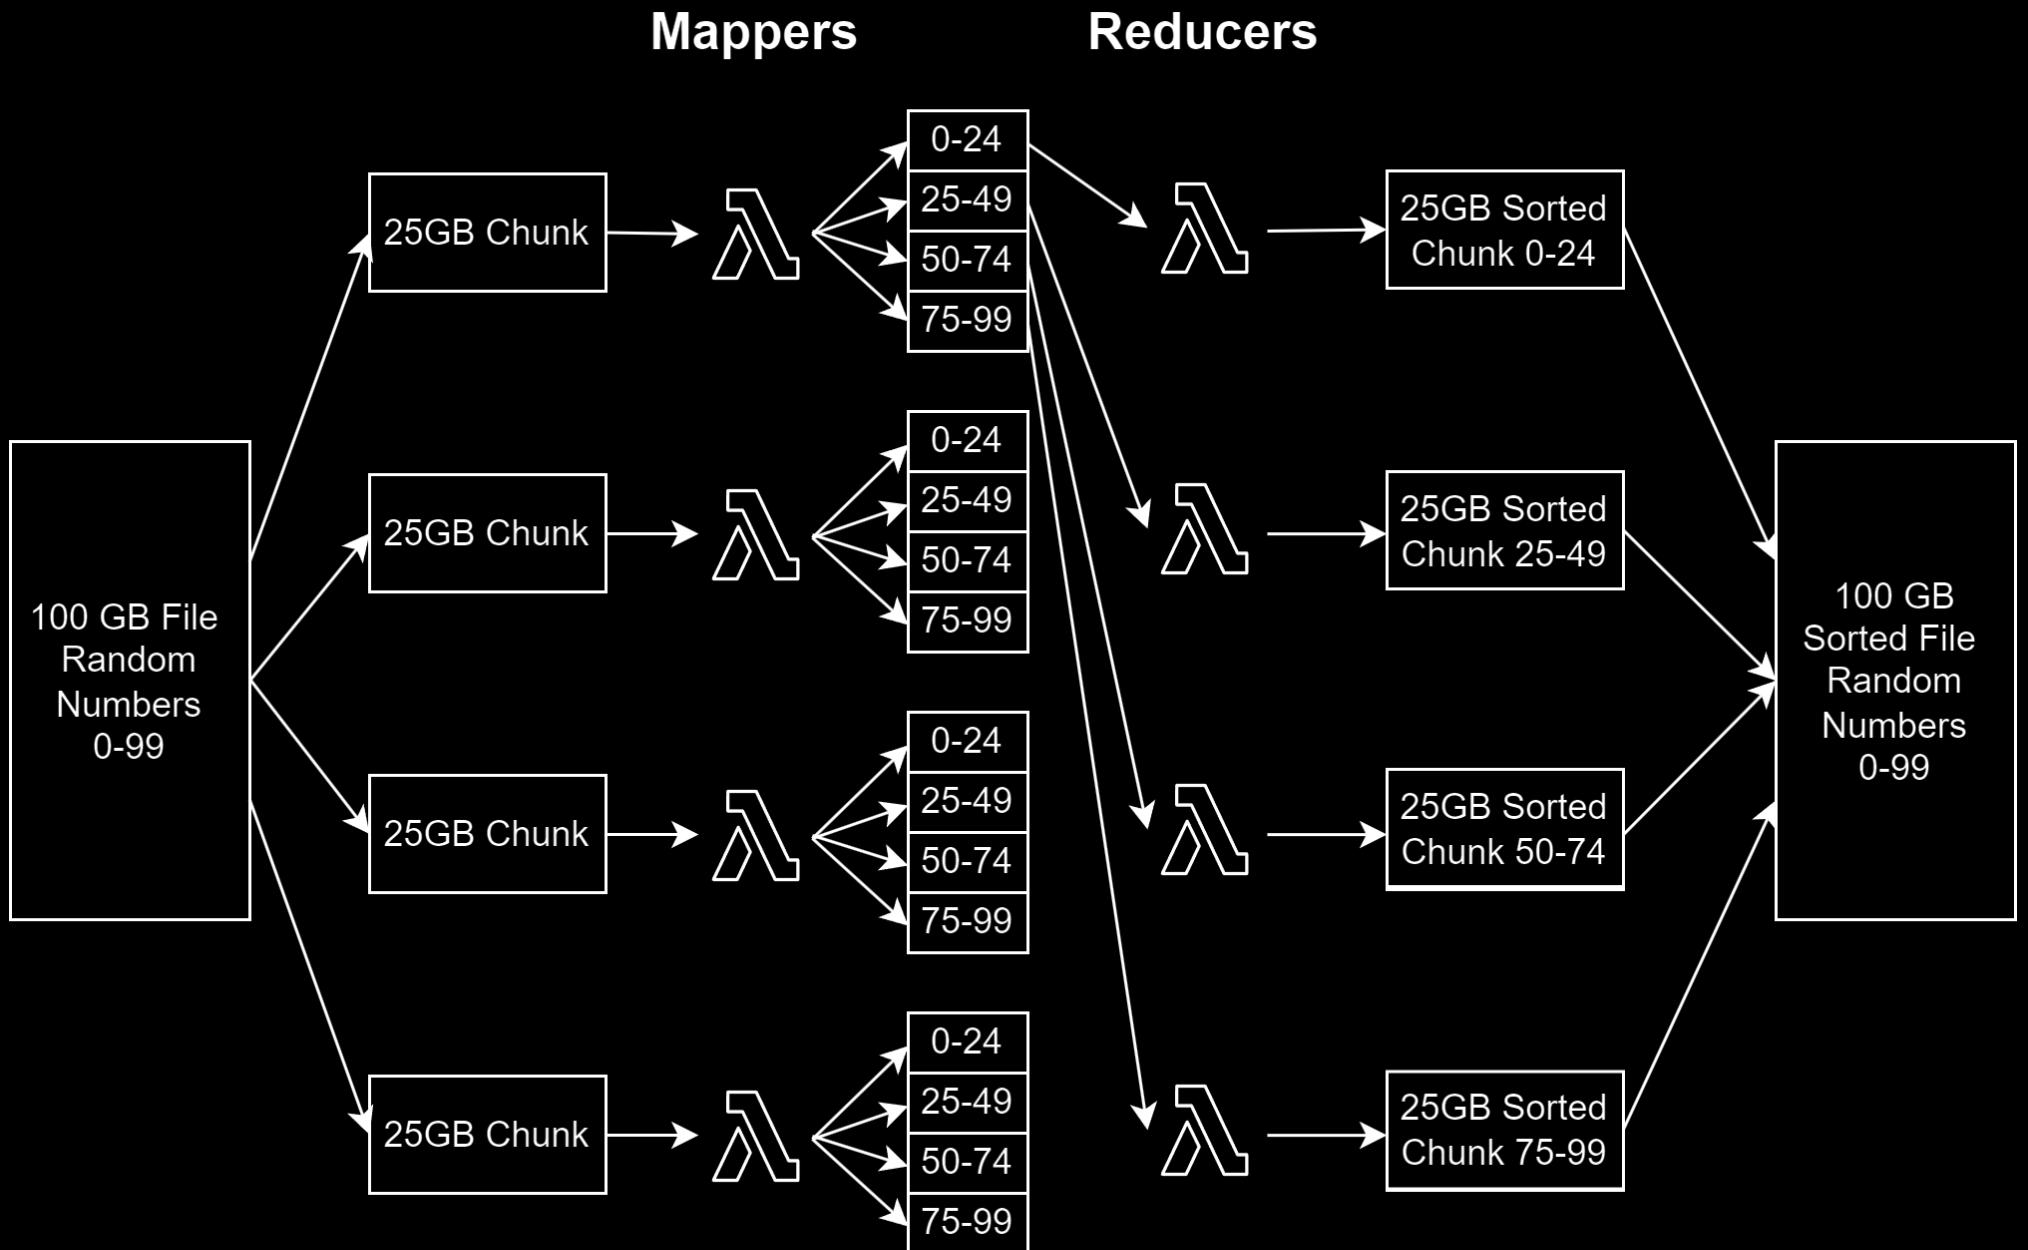
- Why are all these machines so underutilized?
 - Provisioning correctly is very hard
 - **Load changes** a lot over **time**
 - Need to **over-provision**
- If we allocate **resources as function of load**, we can achieve **better utilization**
- **Remove burden** on the developer to scale the deployment

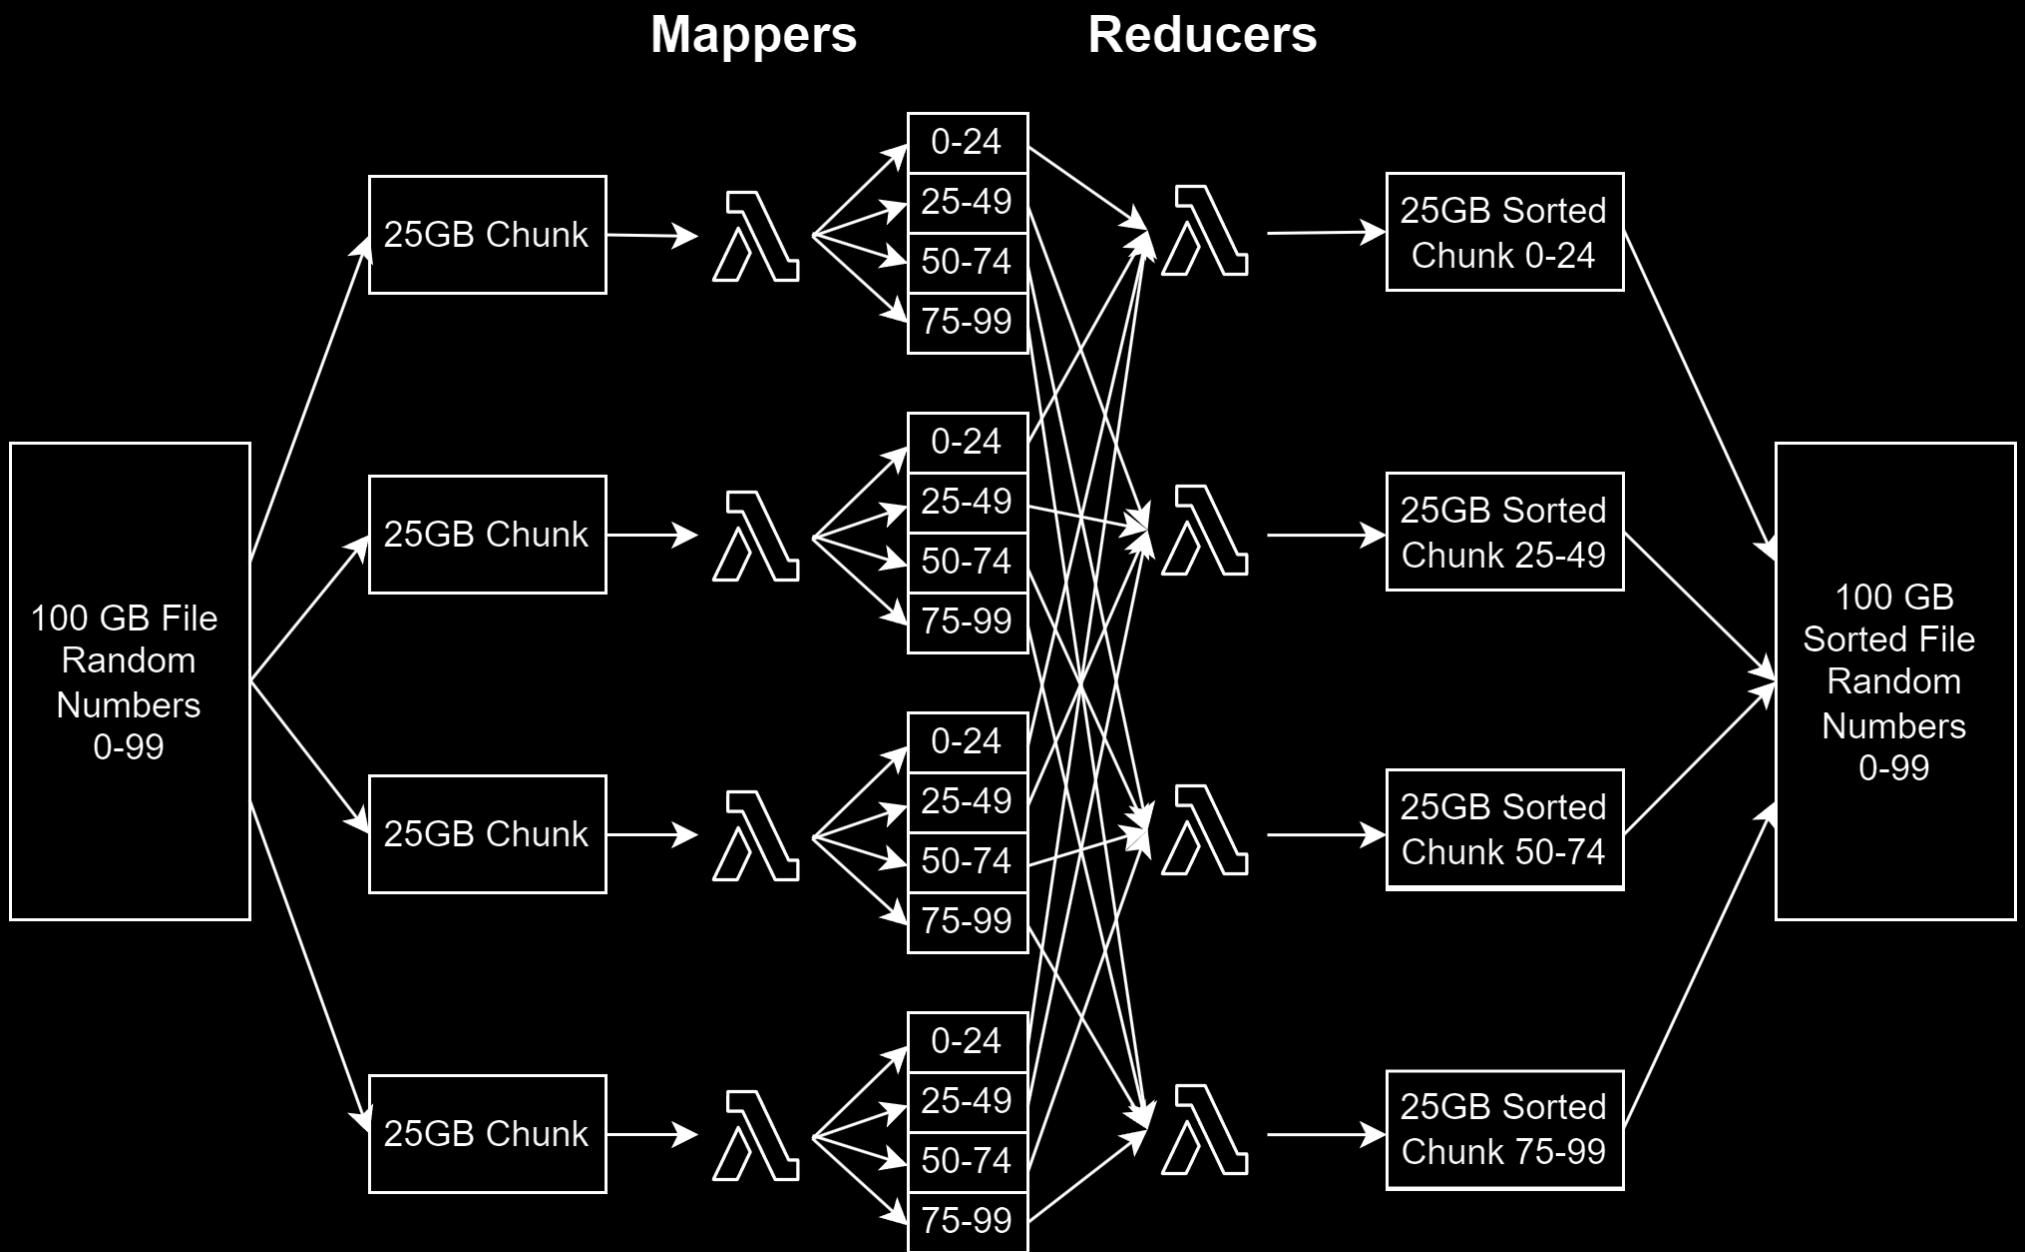
The core values of Serverless

- The **client** expresses applications as **collections of functions**
- The **cloud provider** is responsible for:
 - Allocating resources
 - Load balancing
- The **cloud provider** also:
 - Retains the right to **kill functions** as needed for load balancing
 - Restricts the environment functions run on (e.g., network, local storage)
- The **client** must adapt to this environment

Why no networks for Lambdas

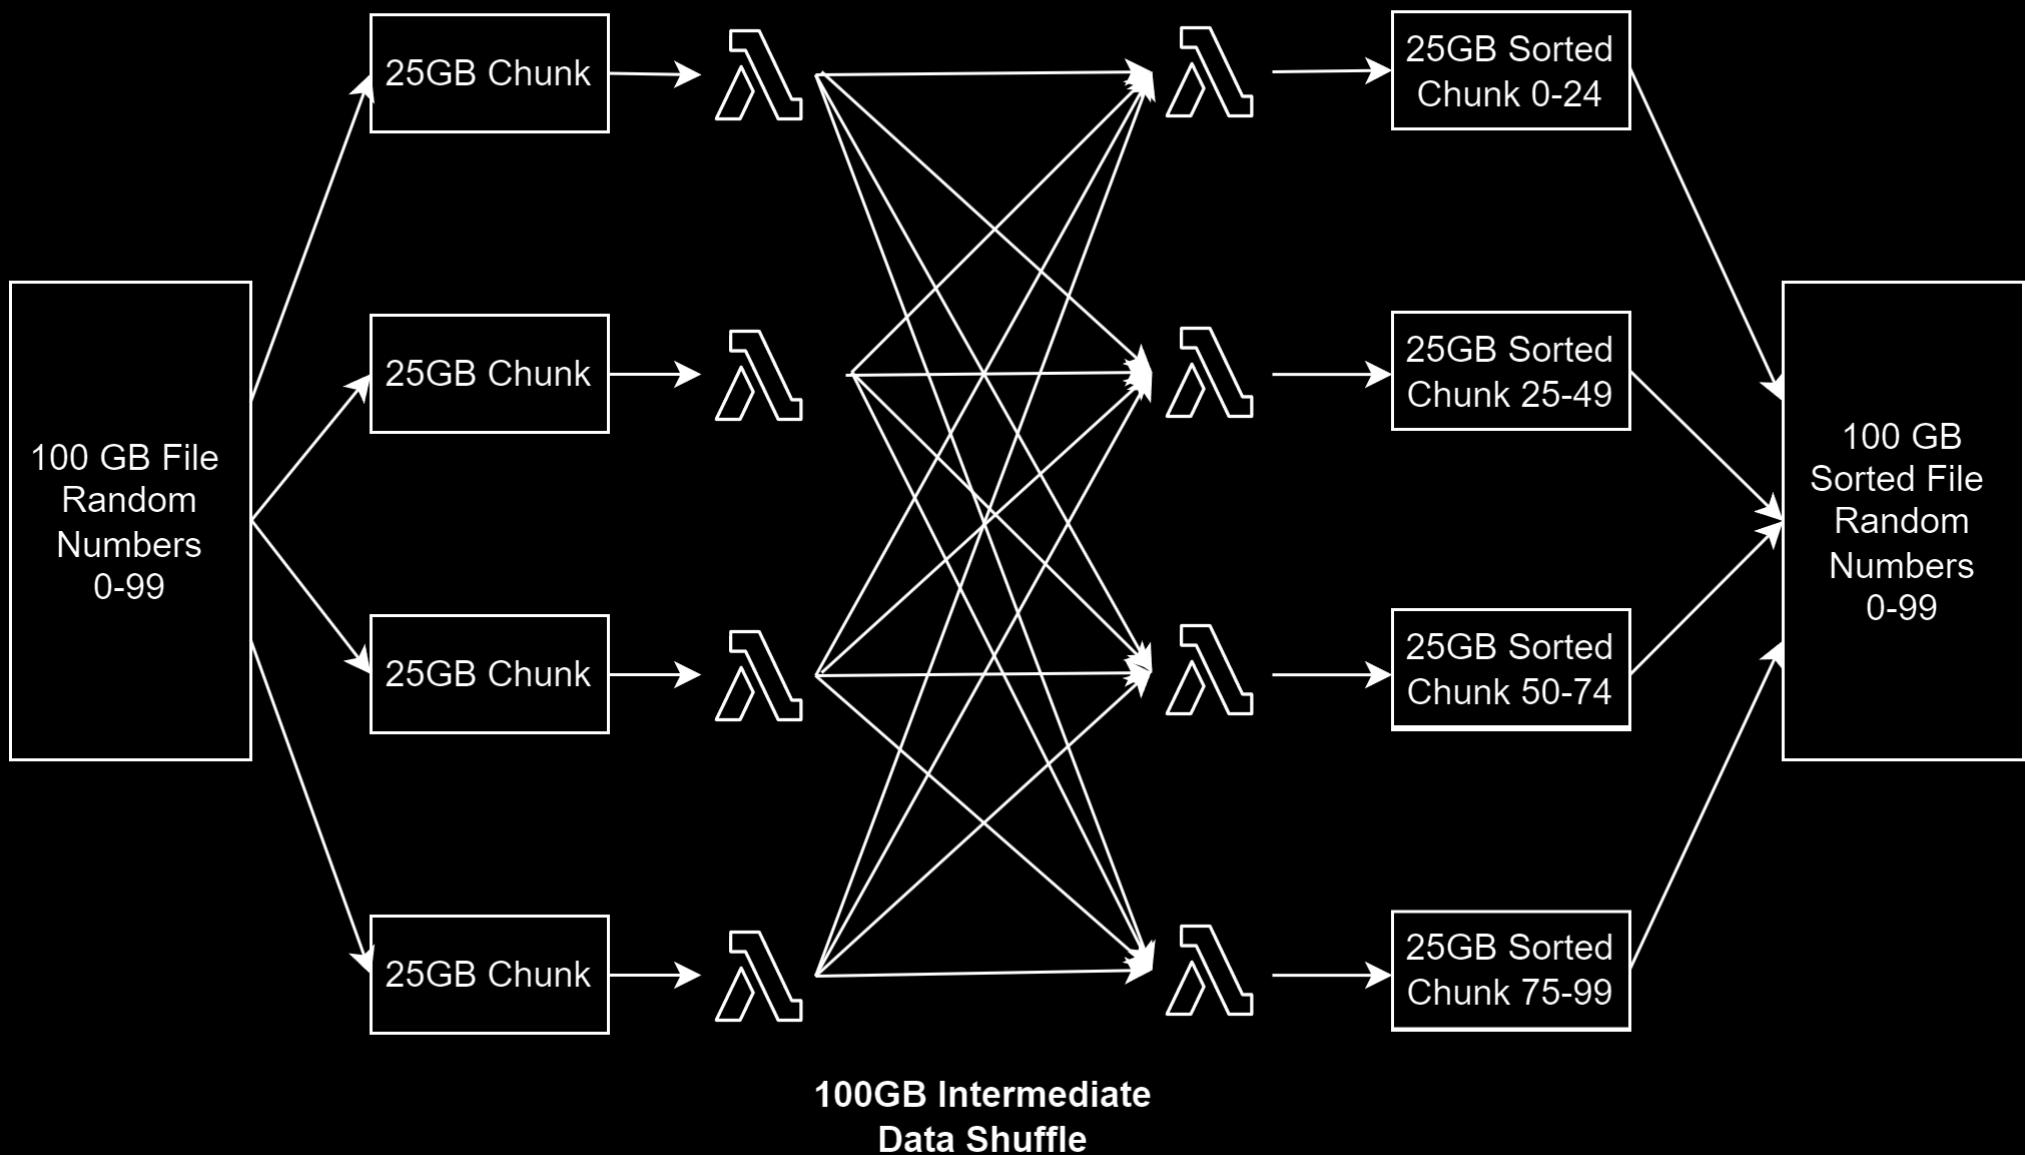
- AWS Lambdas cannot communicate directly (i.e., via IP) with each other
 - This would be useful to build **function pipelines**
- Why?
 - There is no guarantee that Lambda's would run **concurrently**
 - Lambdas can be killed by the administrator, which would **create a cascading failure**
- How is it done in practice?
 - Communicate asynchronously through a storage service (e.g., S3)
 - Expensive, slow and sub-optimal

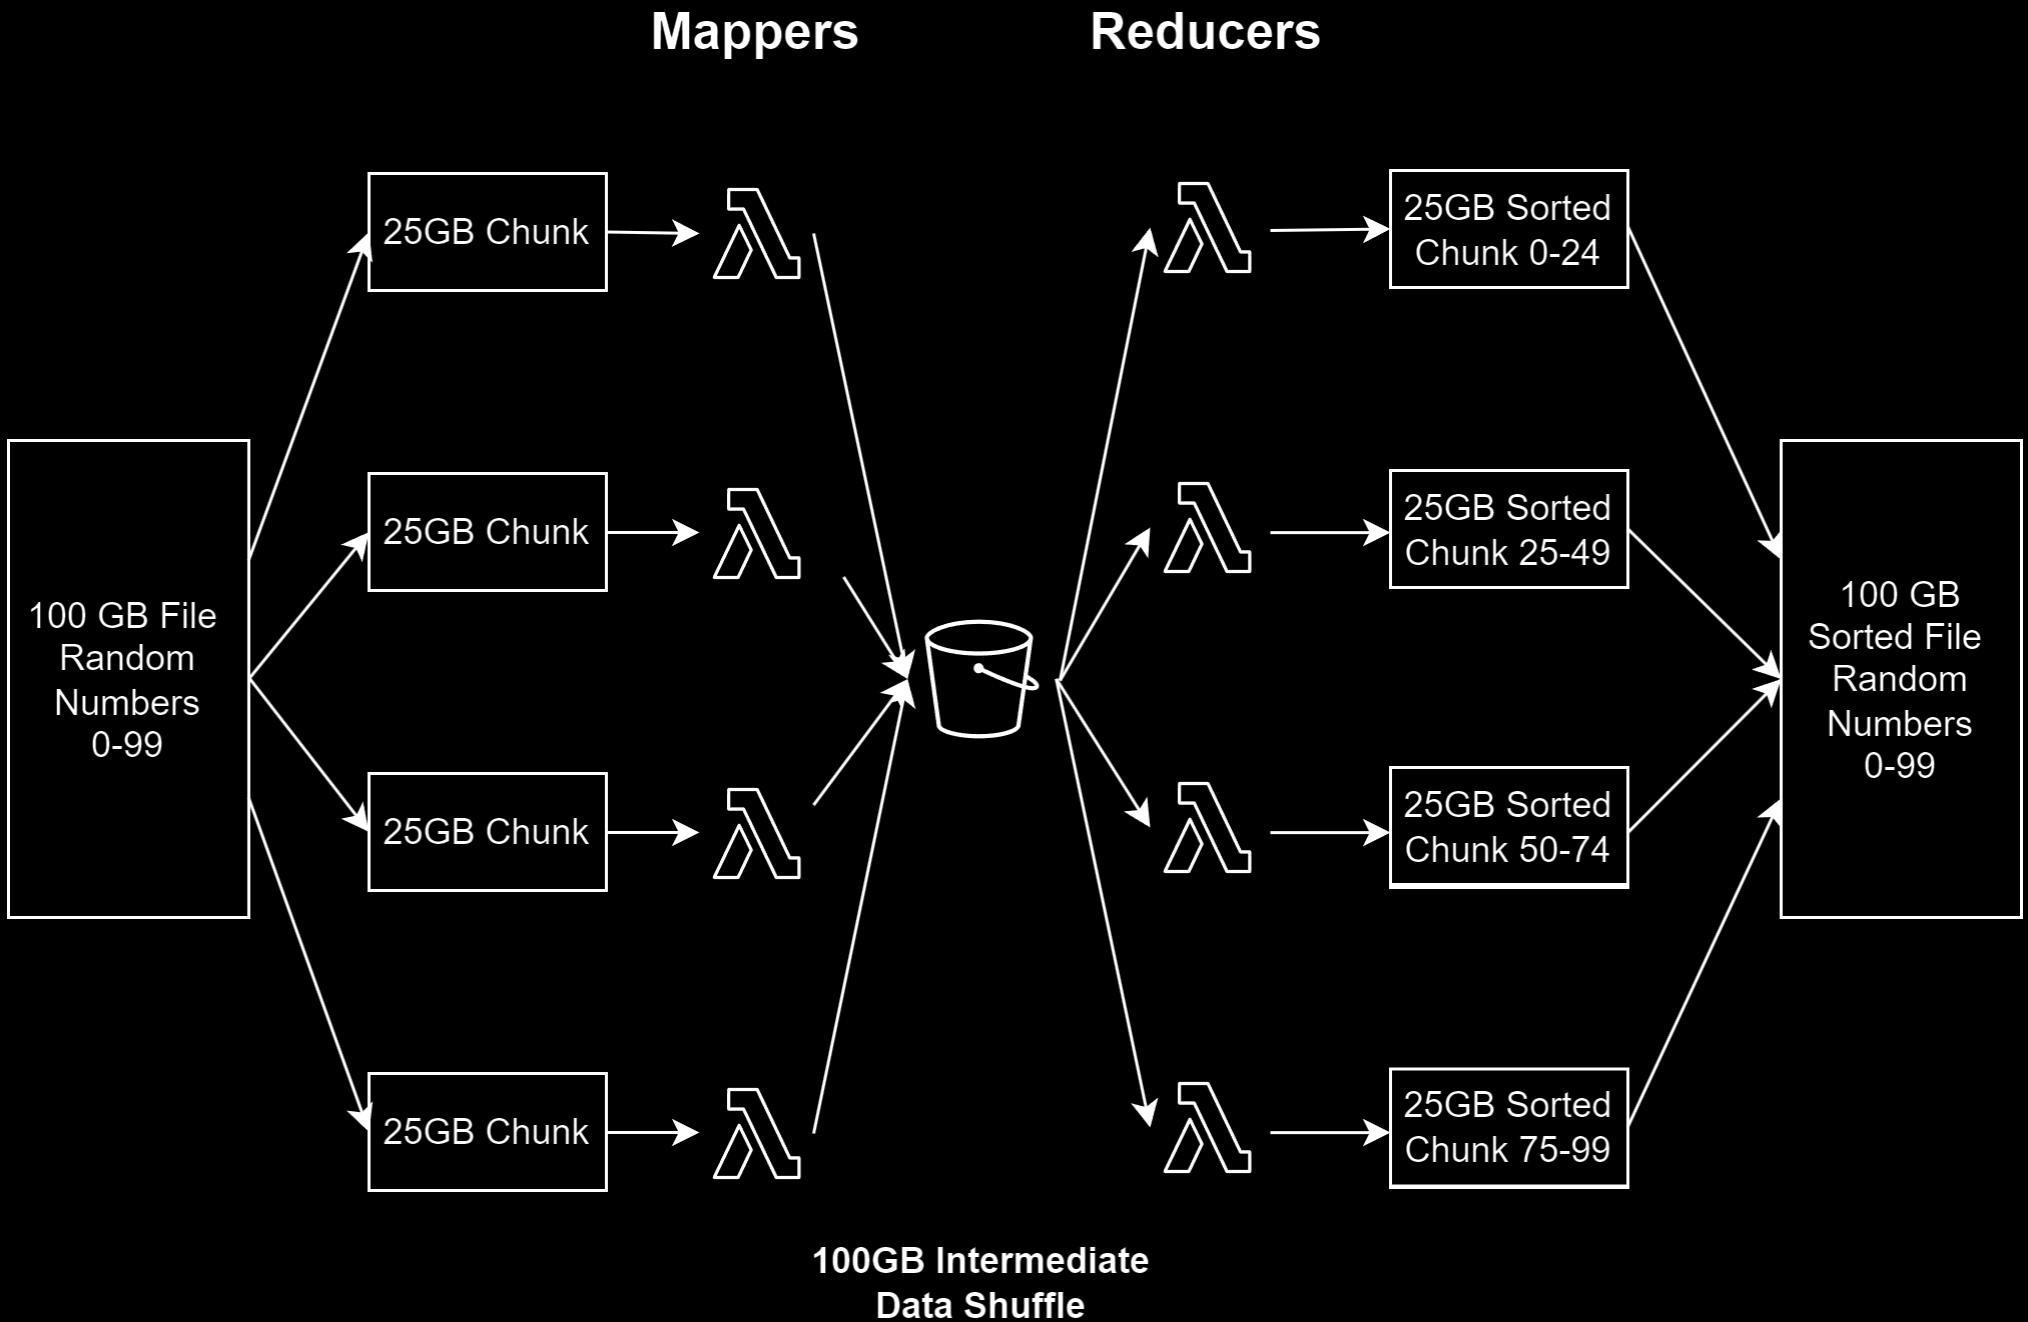




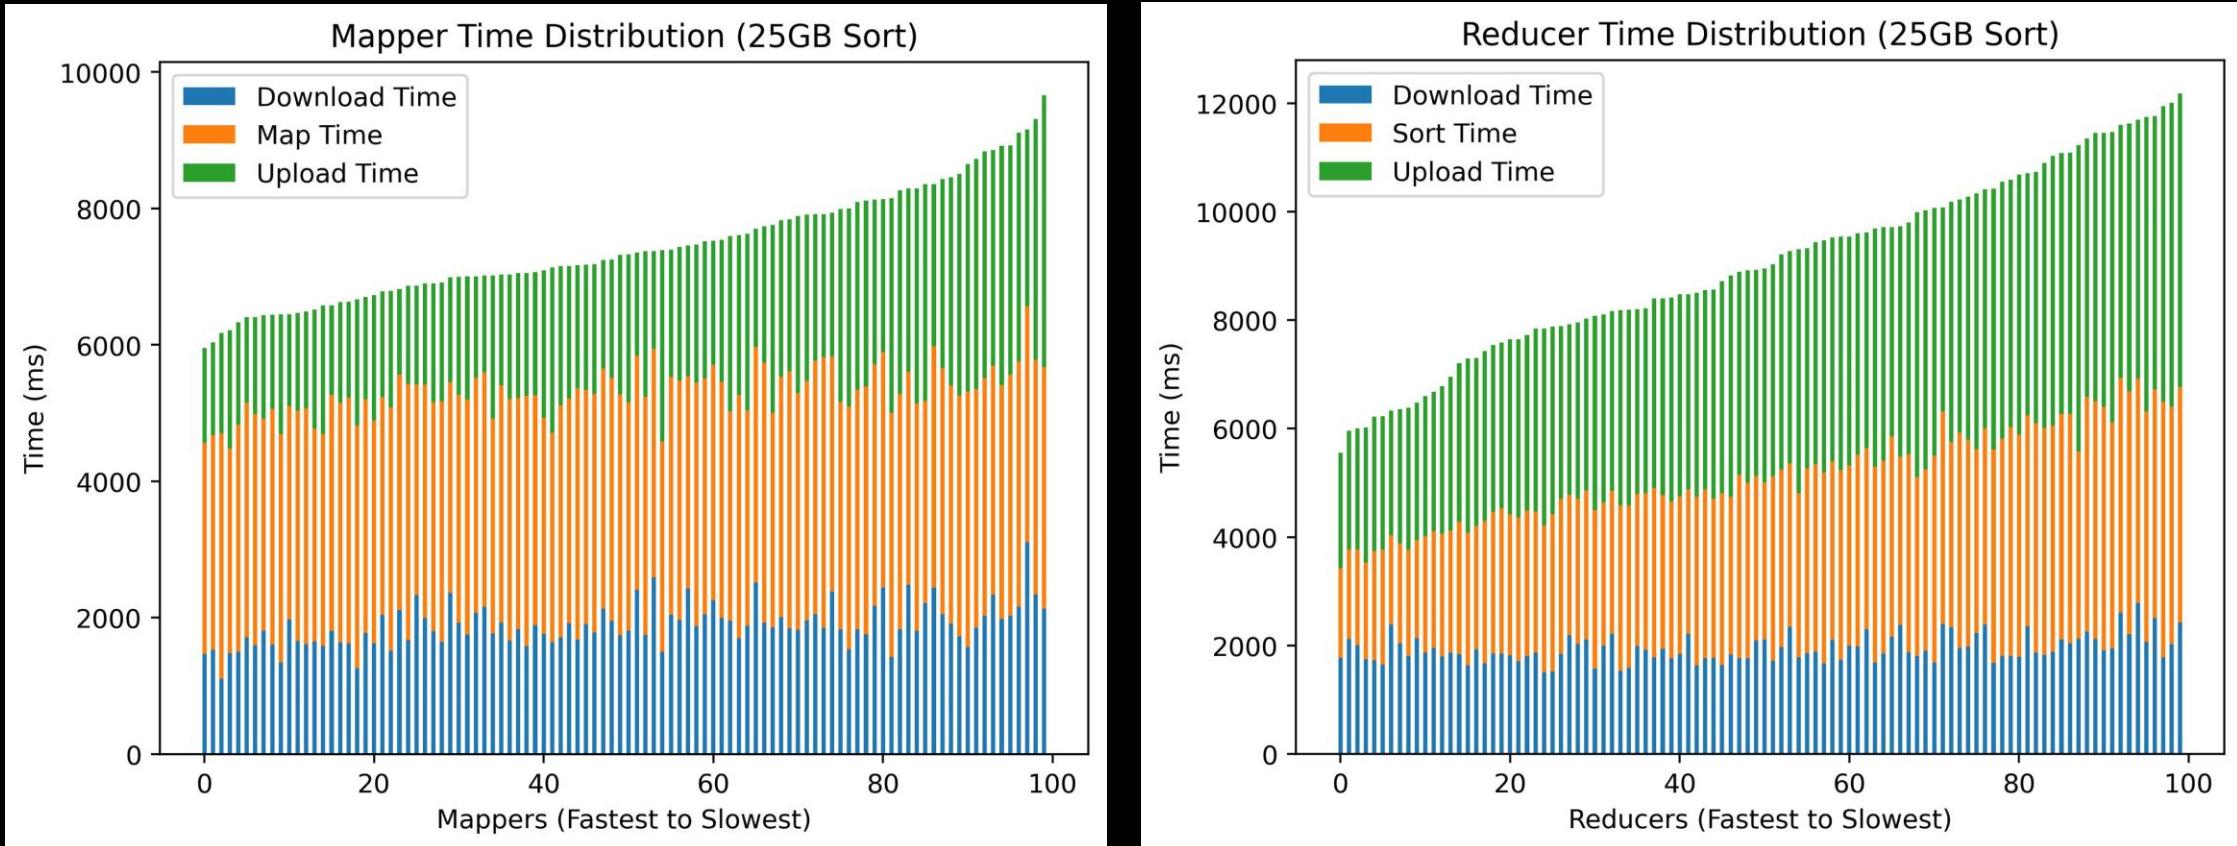
Mappers

Reducers





Communication **is** a bottleneck



Latency breakdown of compute time and communication time (through S3) of a MapReduce-like job implemented using Lambdas communicating over S3

Roadmap

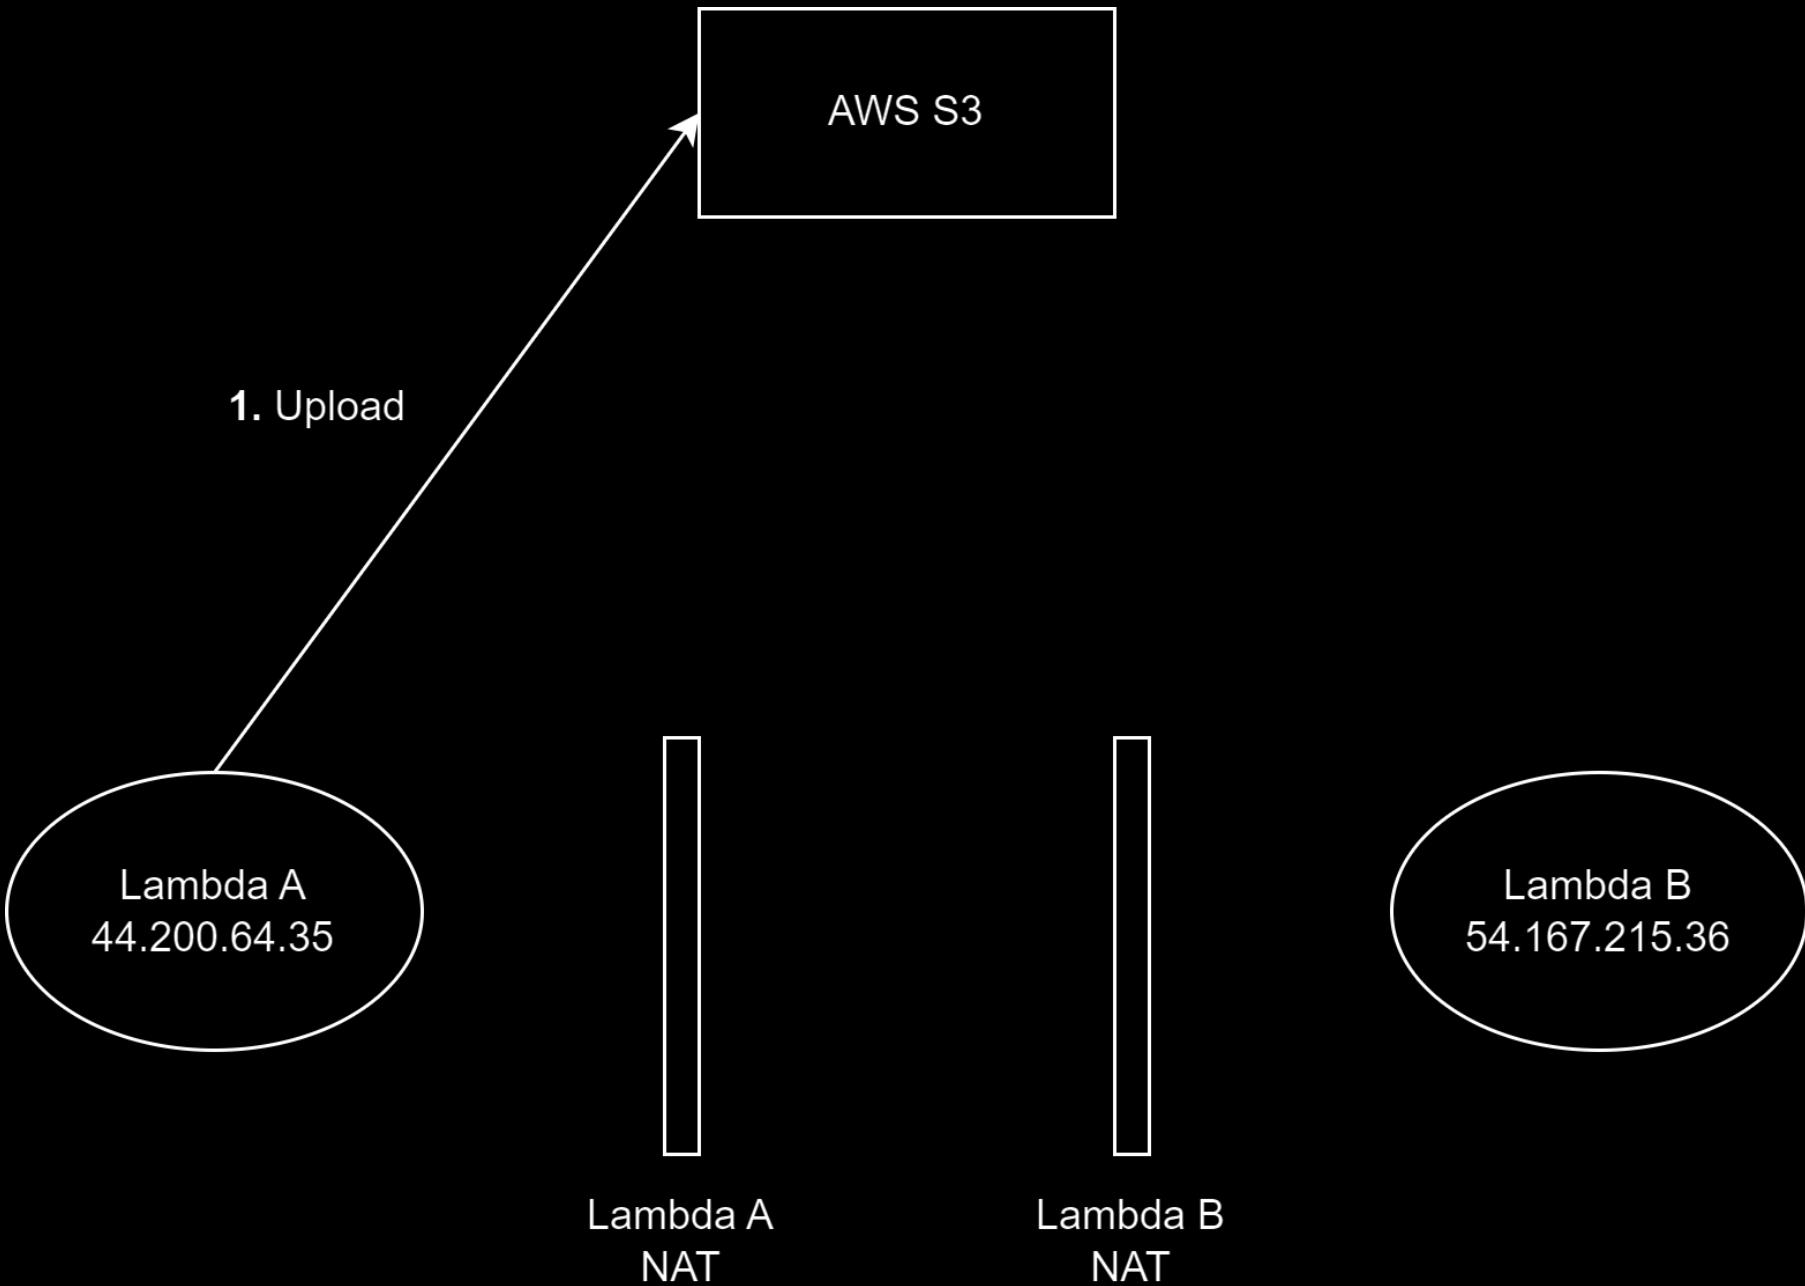
- Serverless primer
 - Why serverless
 - How serverless
 - Why no serverless + networks (so far)
- Application Example
 - MapReduce framework
- Punching holes in the Network for fun and profit
 - Technique
 - Battle Scars
 - Performance, or *the great S3 conspiracy*
- Comparison with AWS EMR
- Making it resilient
- Conclusion & Future Work

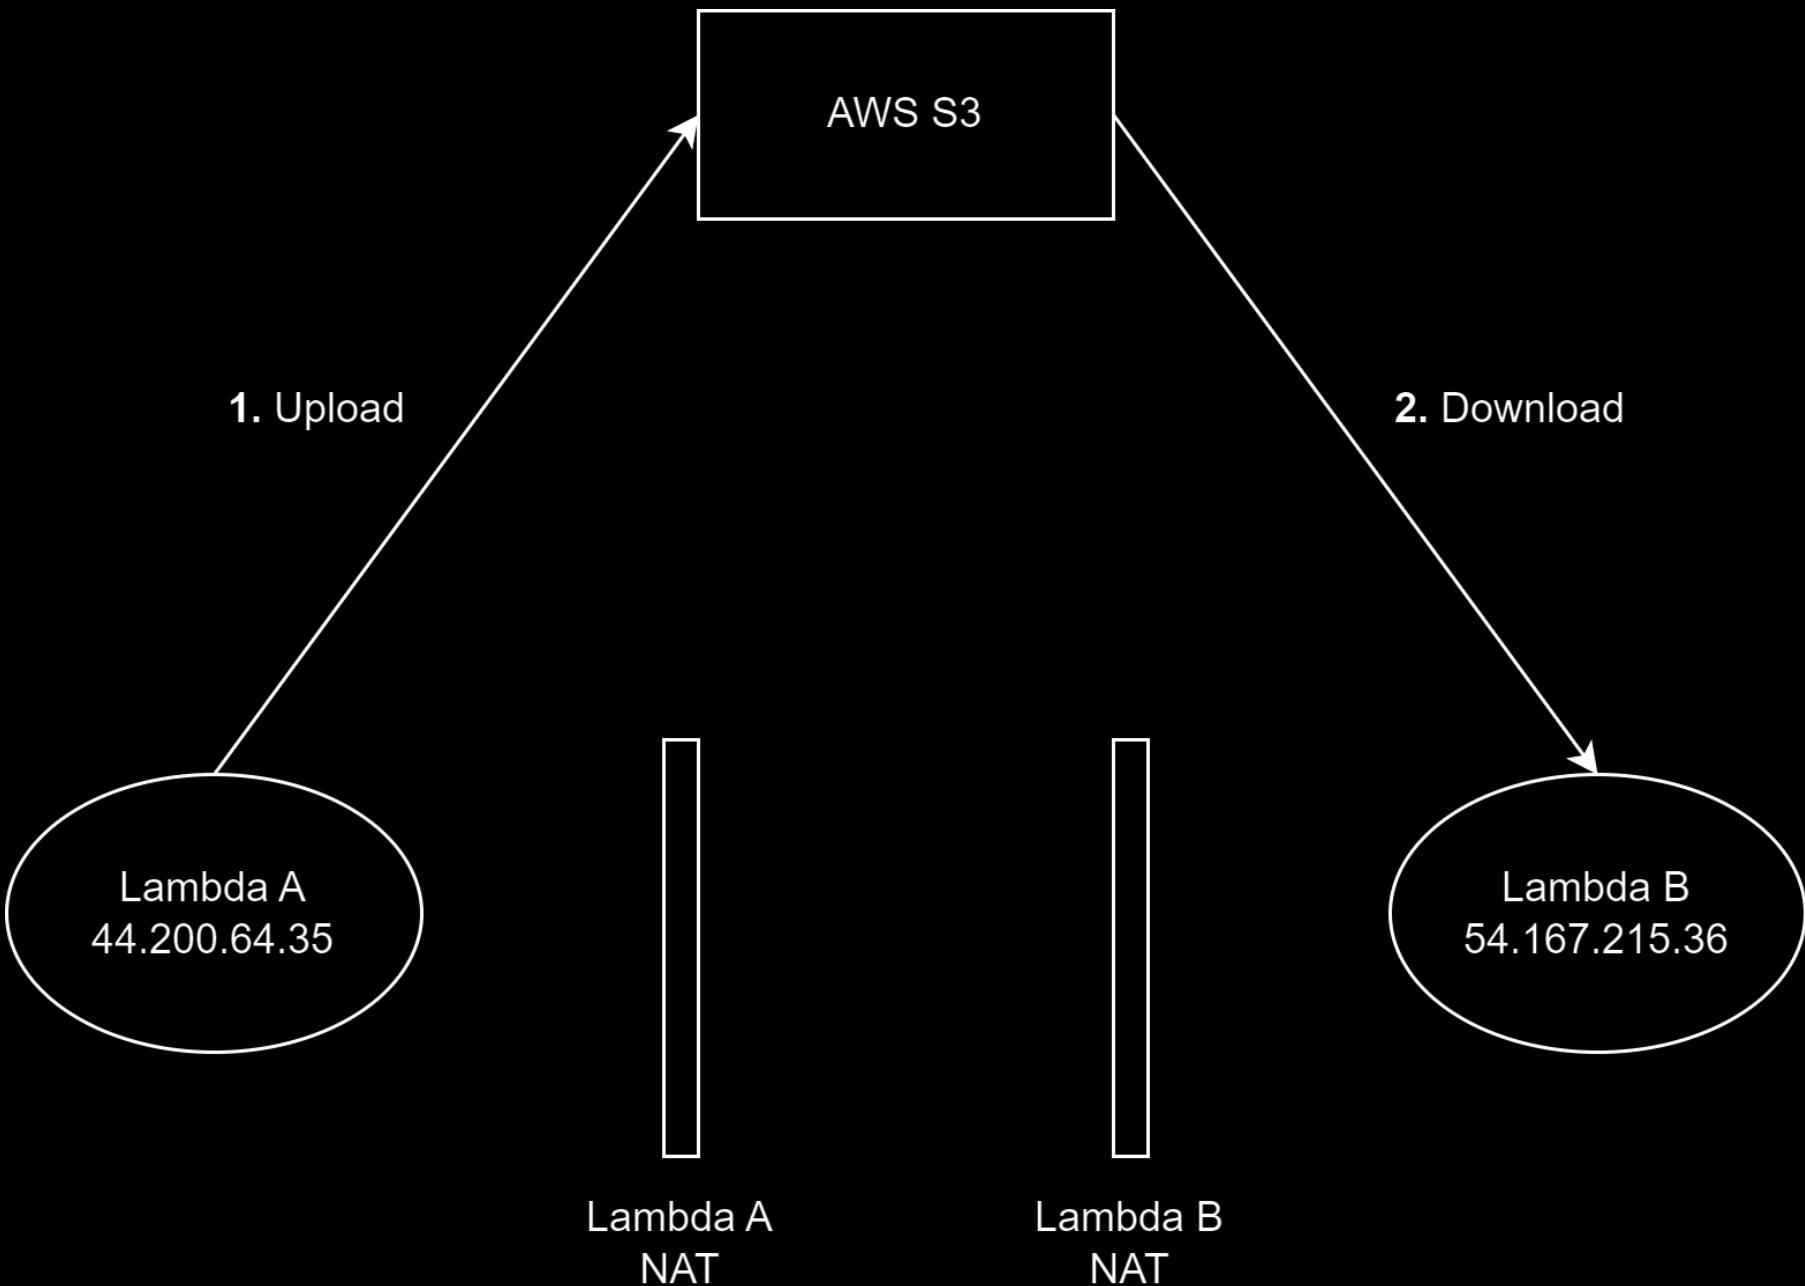


Punching Holes for Fun and Profit









Hole Punching

- Establishes a direct bi-directional link between NATed devices
- Popular technique in peer-to-peer applications
- Introduces **overhead** to establish a link
- Once established, a link is a standard UDP connection
- Maintaining a link requires keep-alive packets

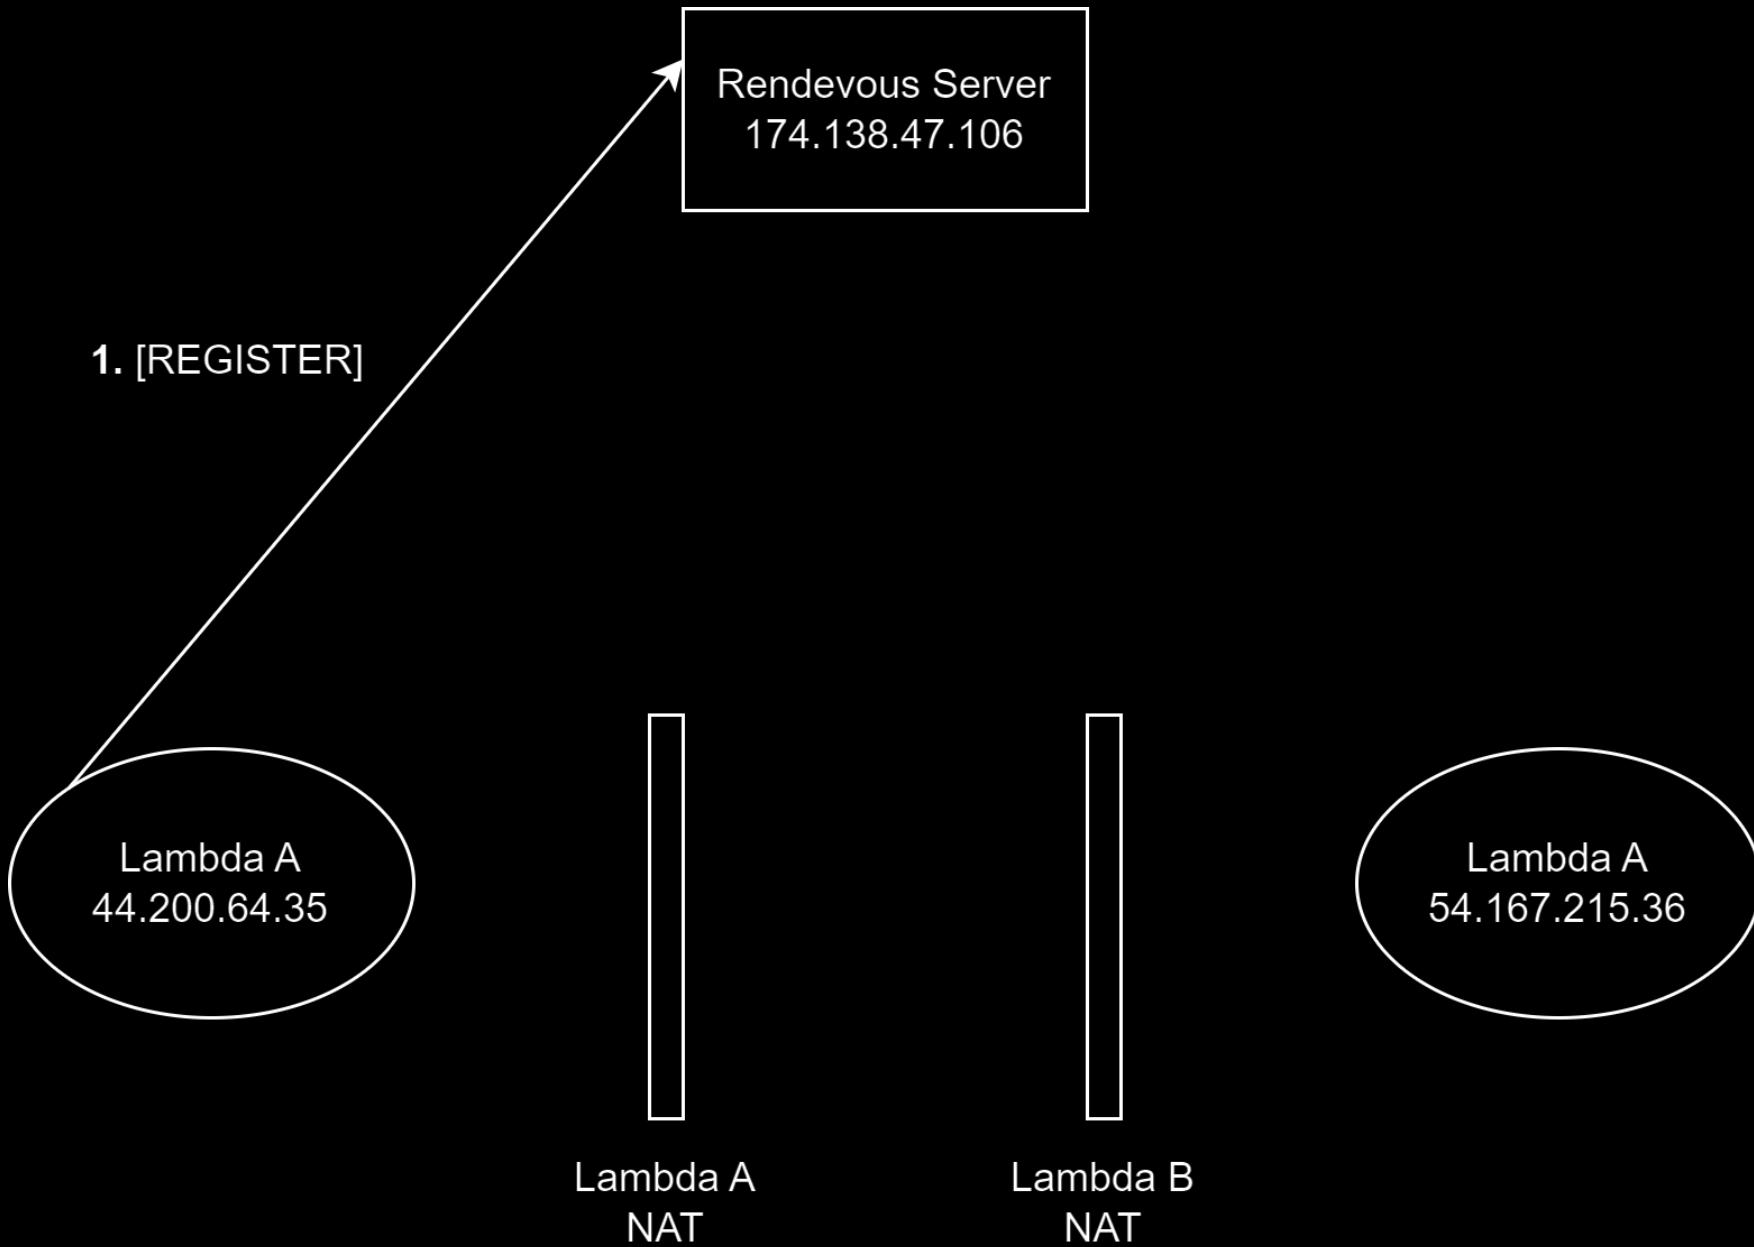
Rendevous Server
174.138.47.106

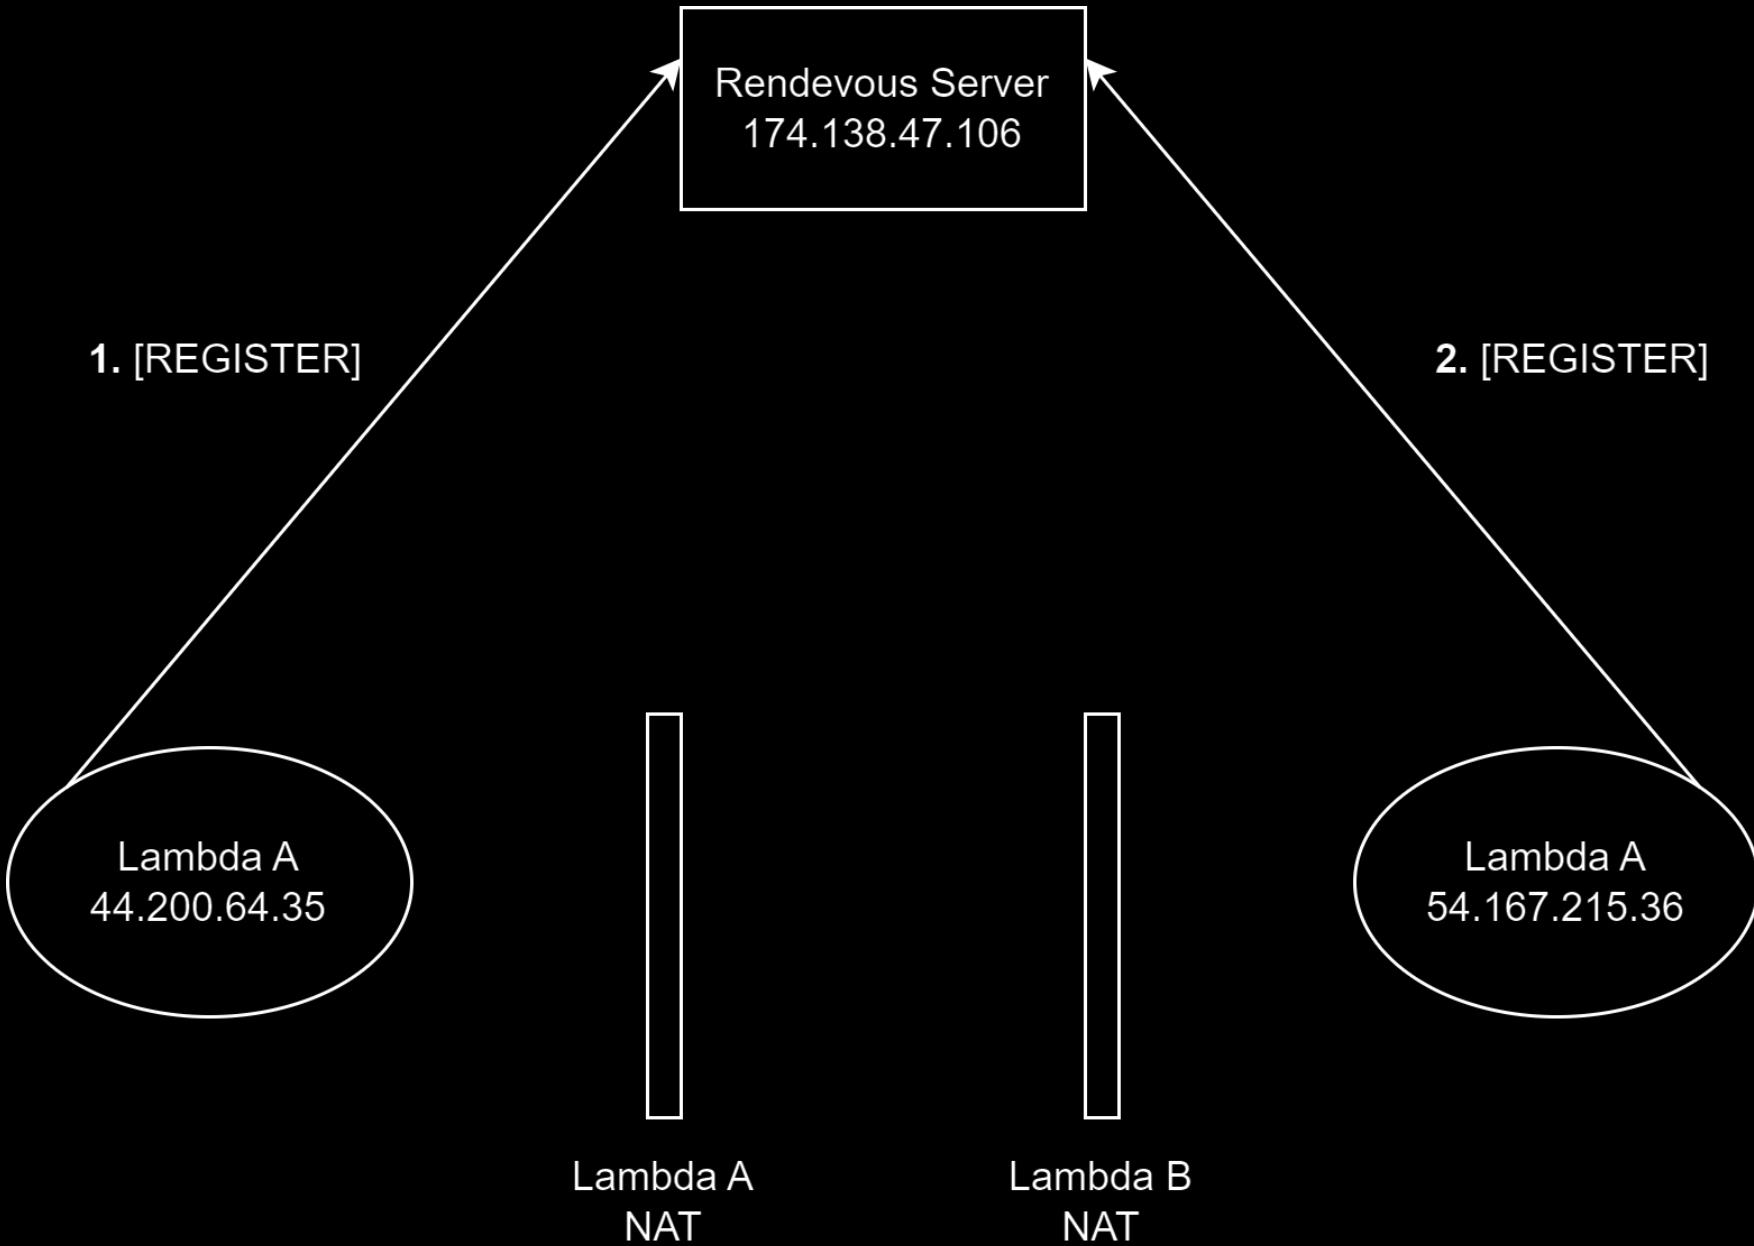
Lambda A
44.200.64.35

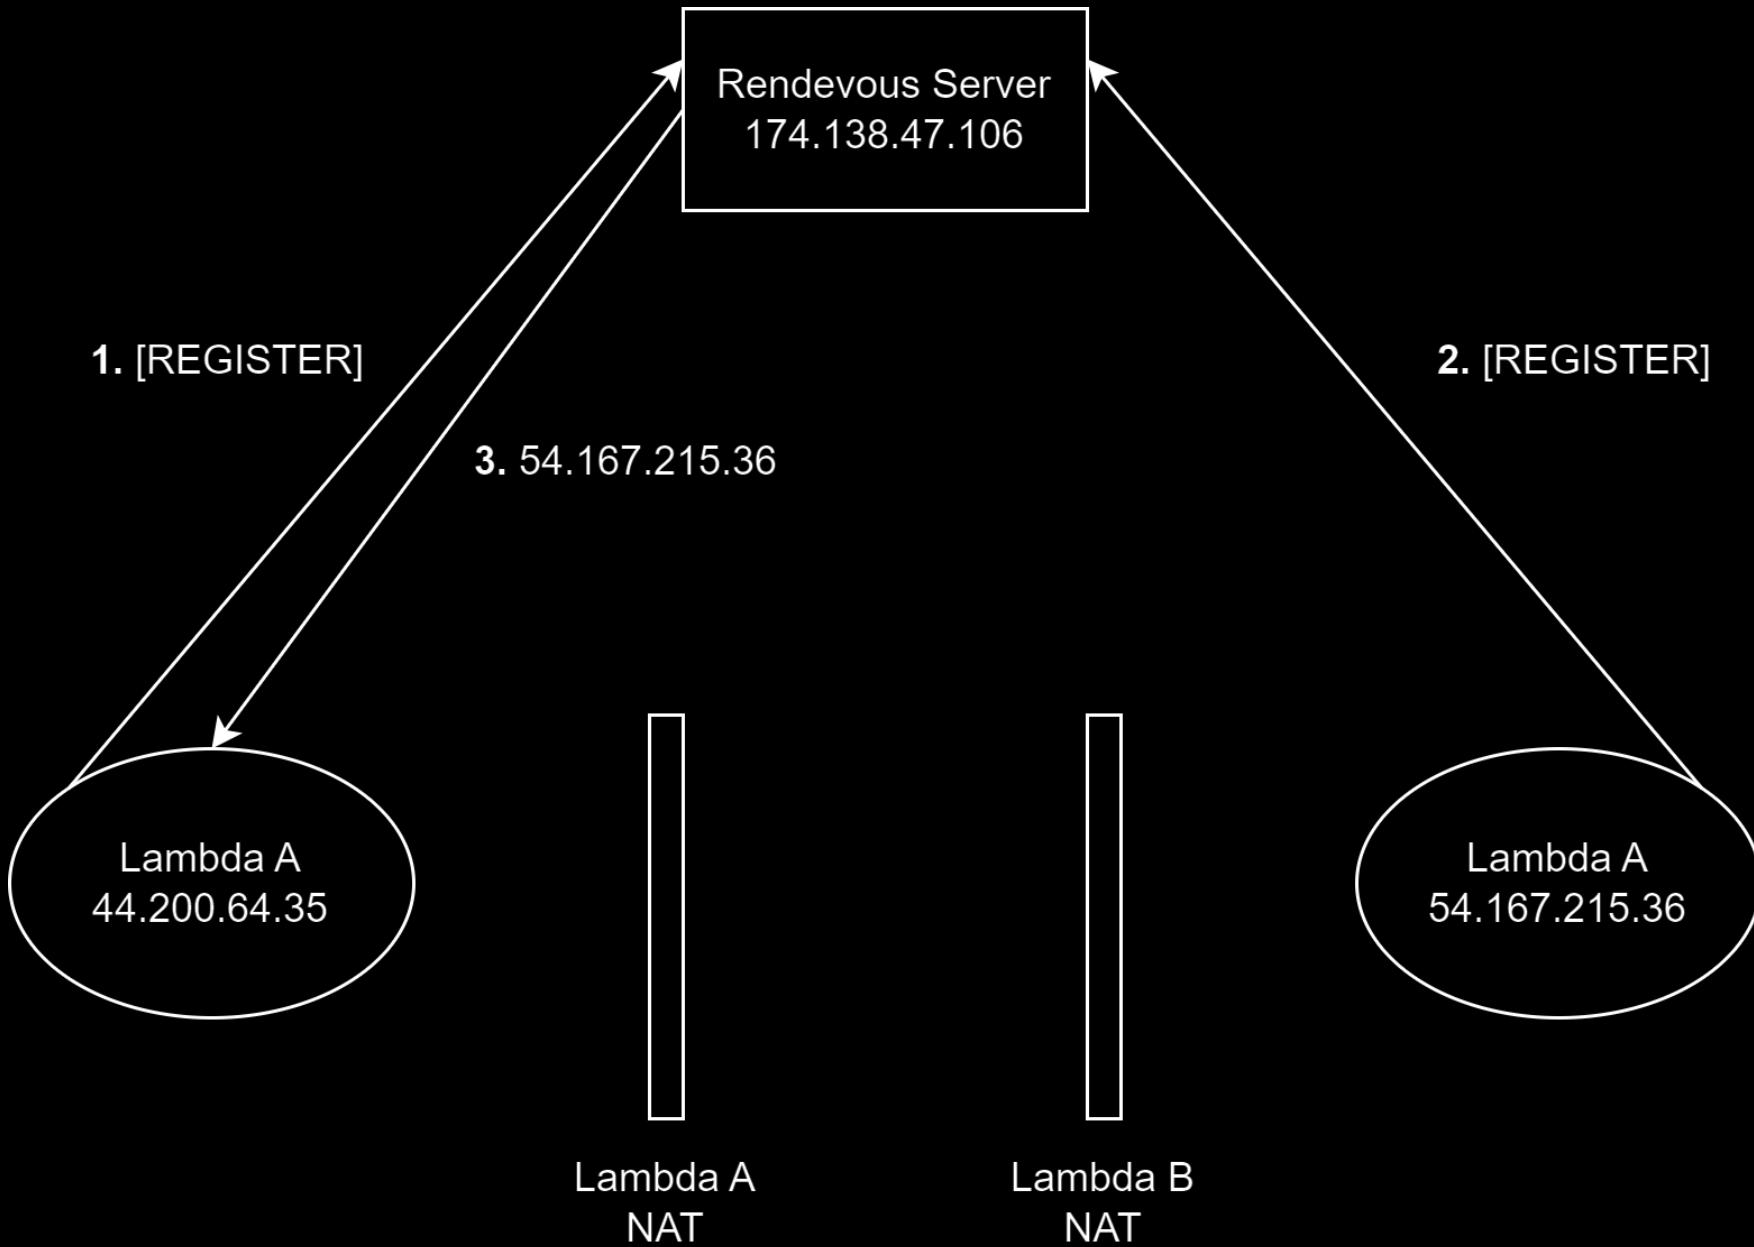
Lambda A
NAT

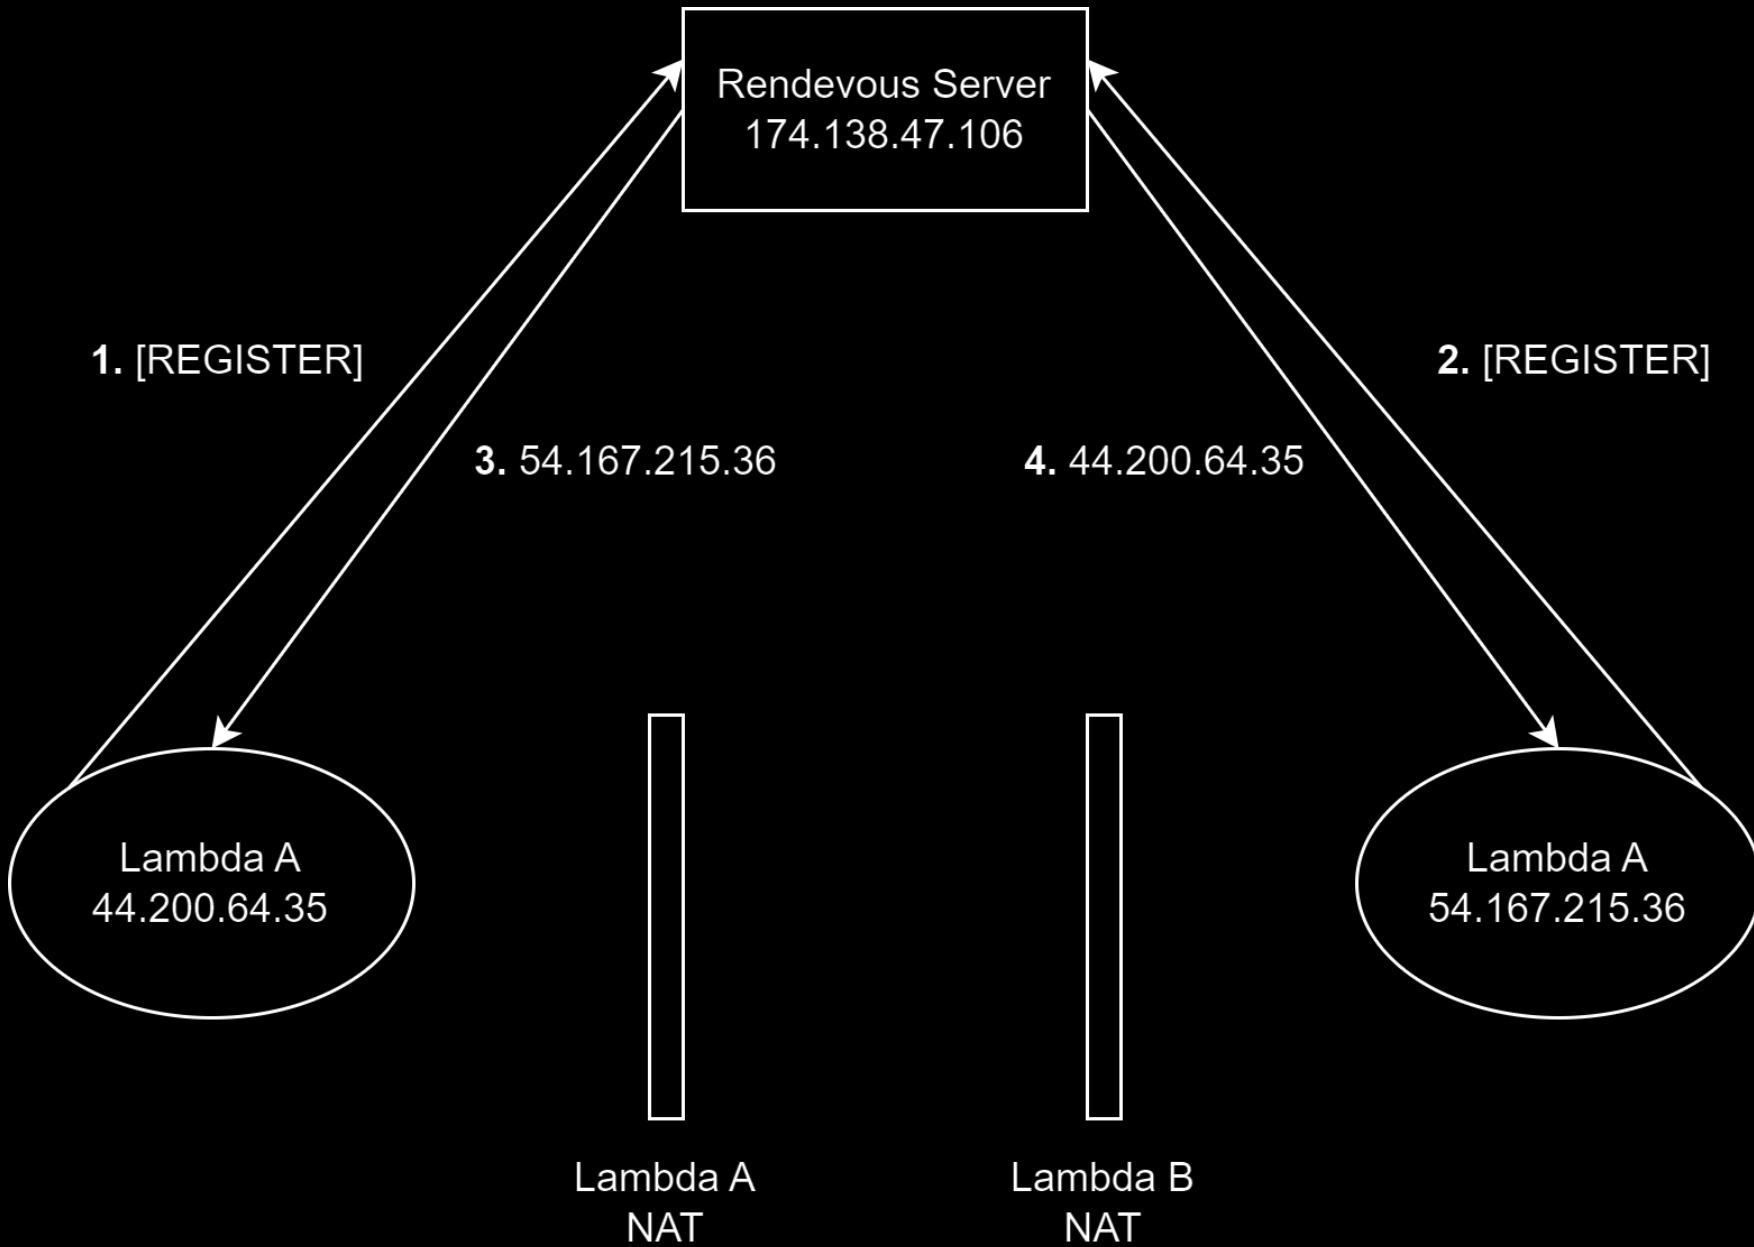
Lambda B
NAT

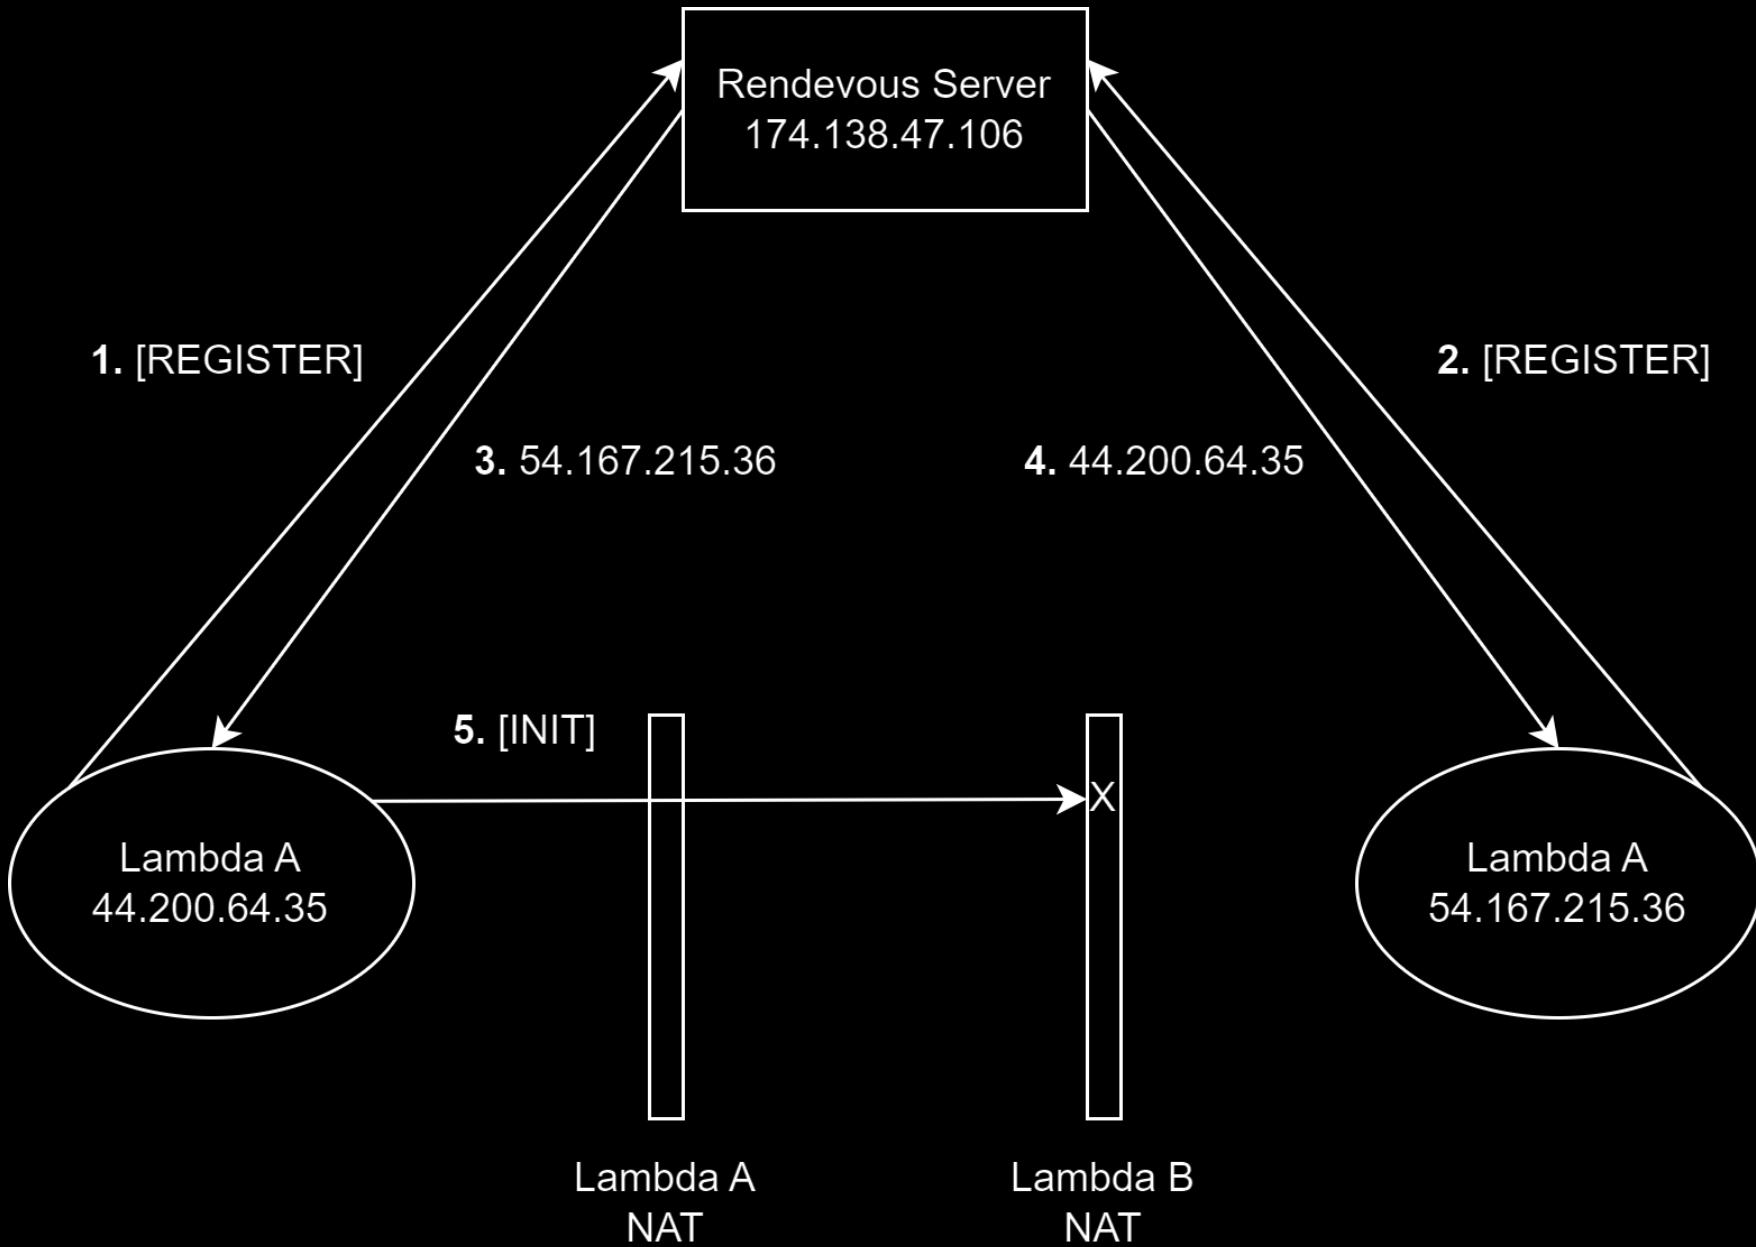
Lambda A
54.167.215.36

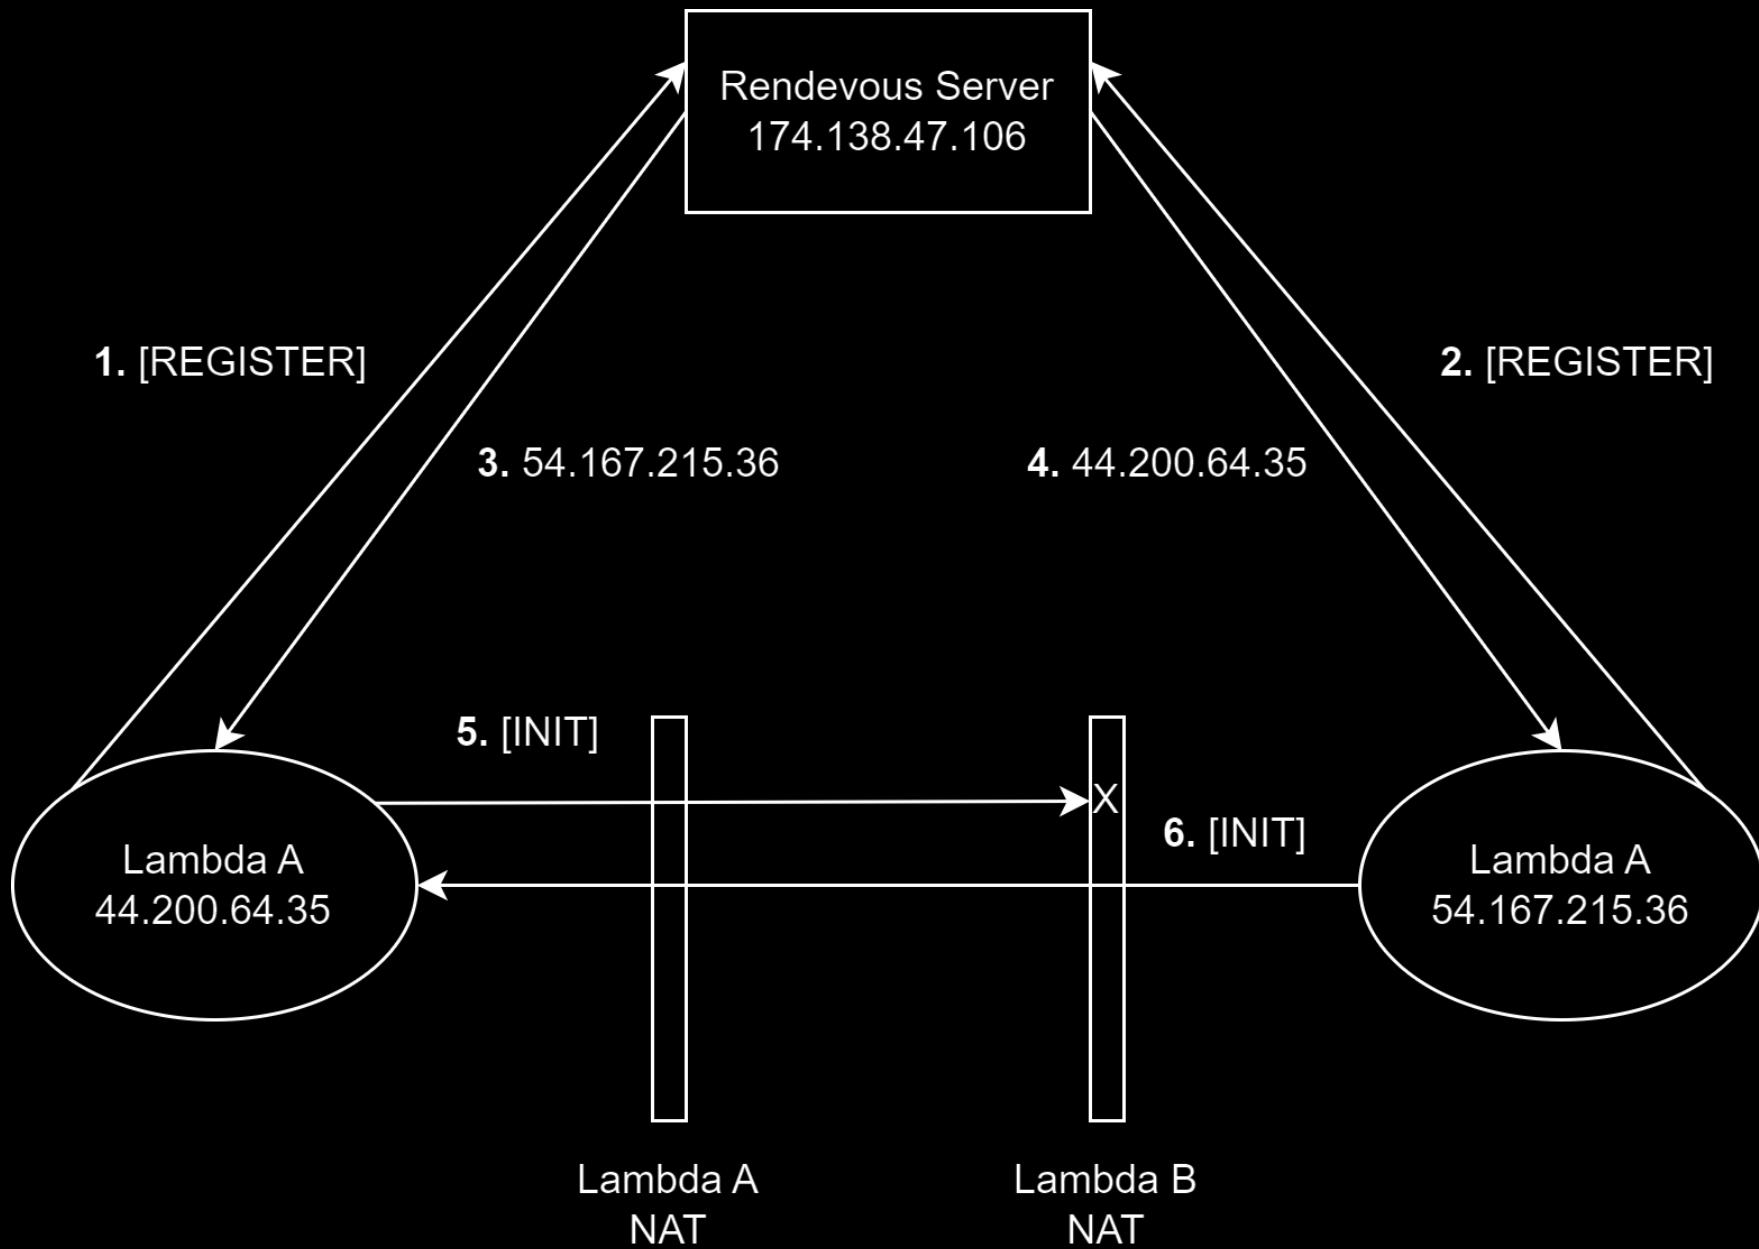


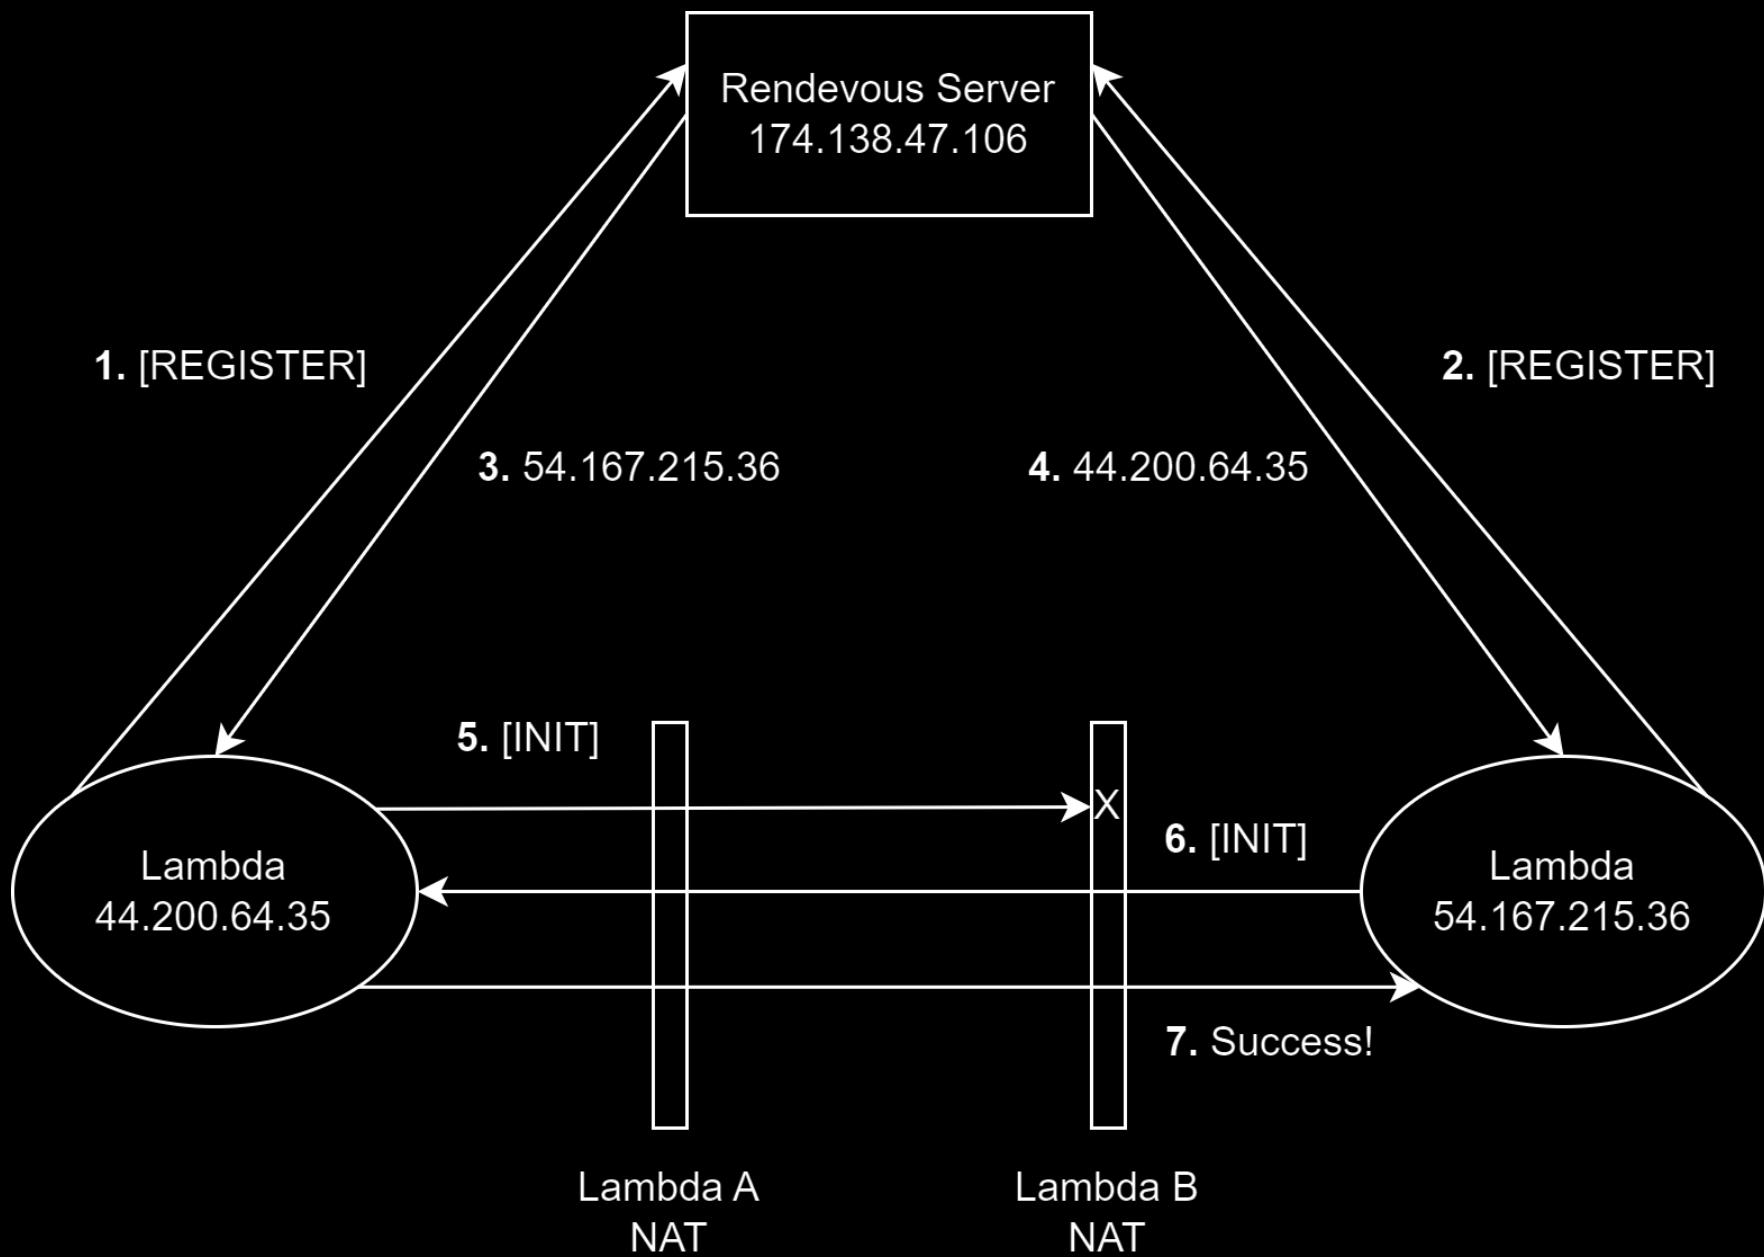






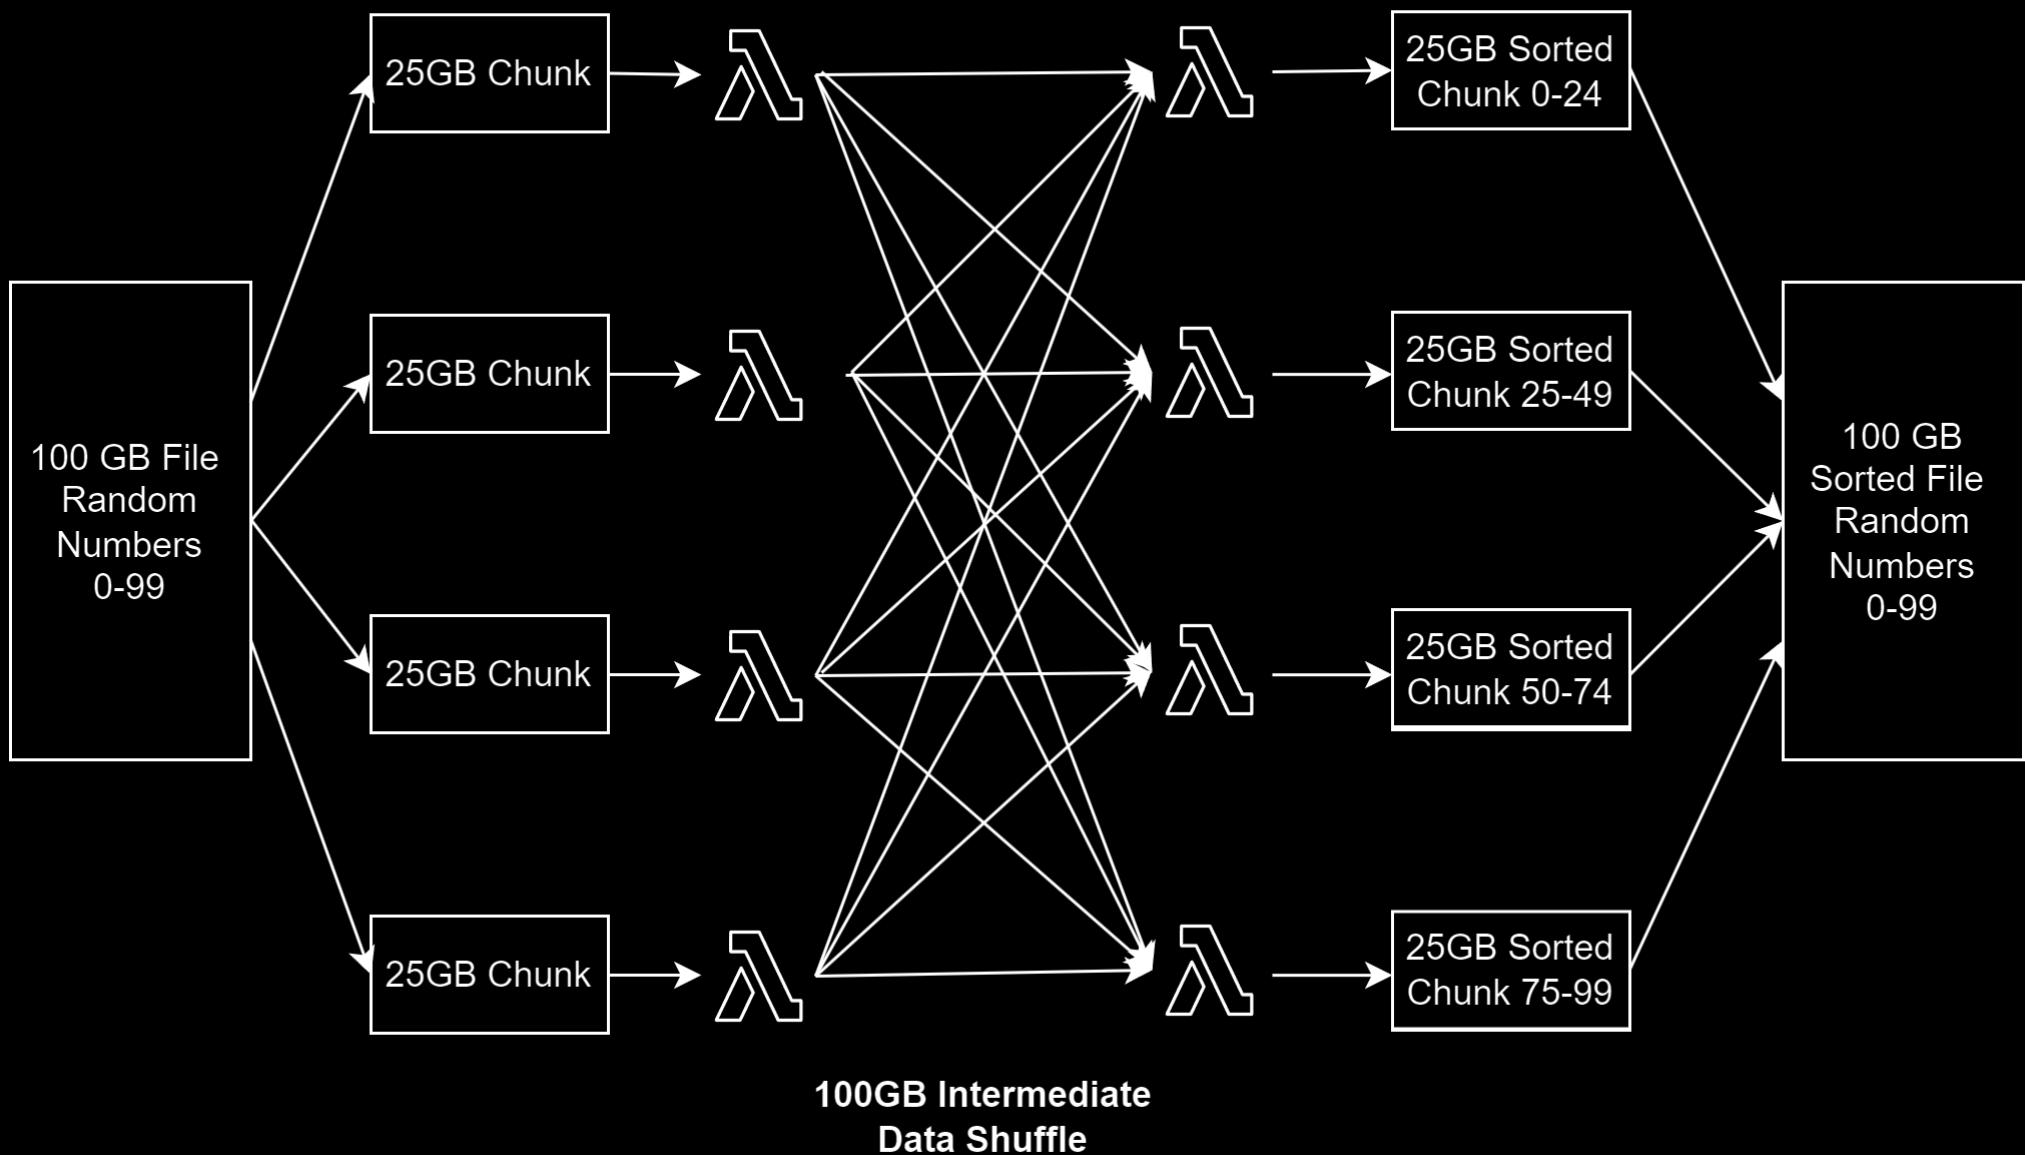




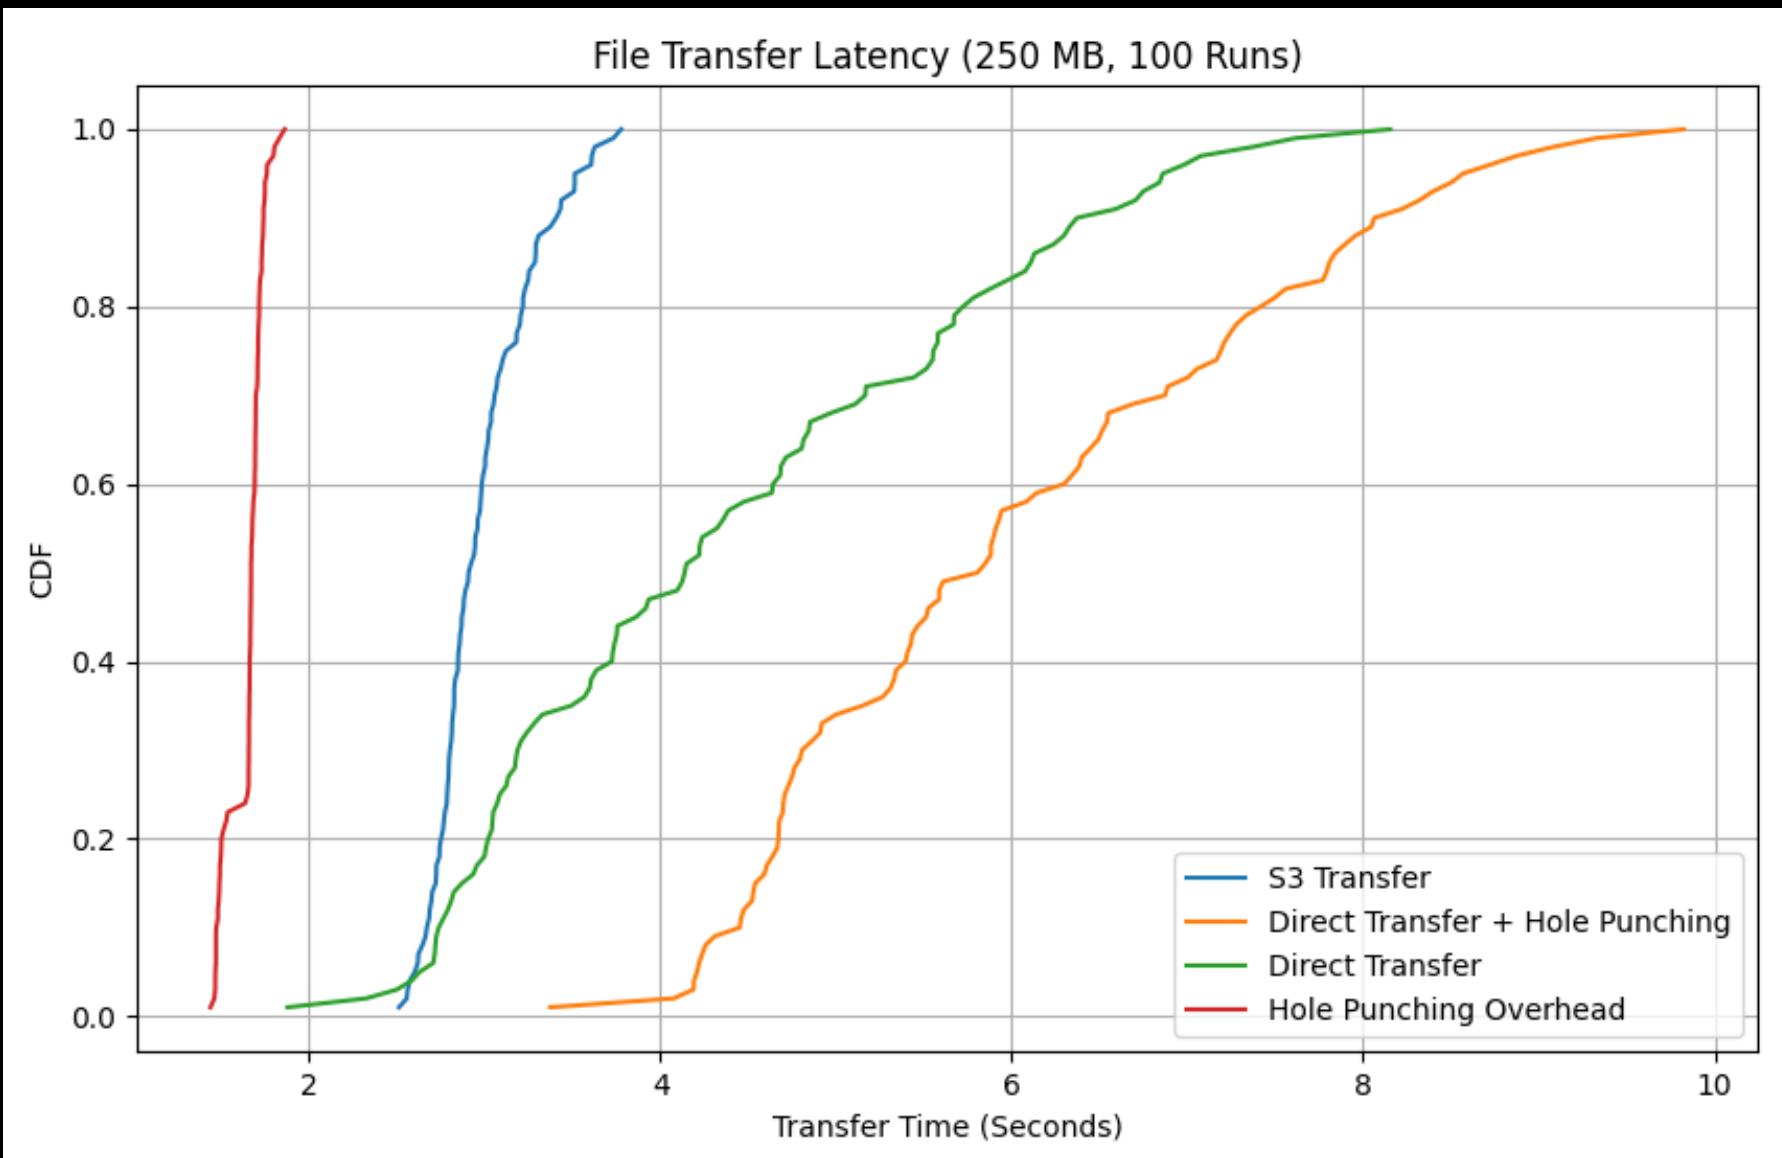


Mappers

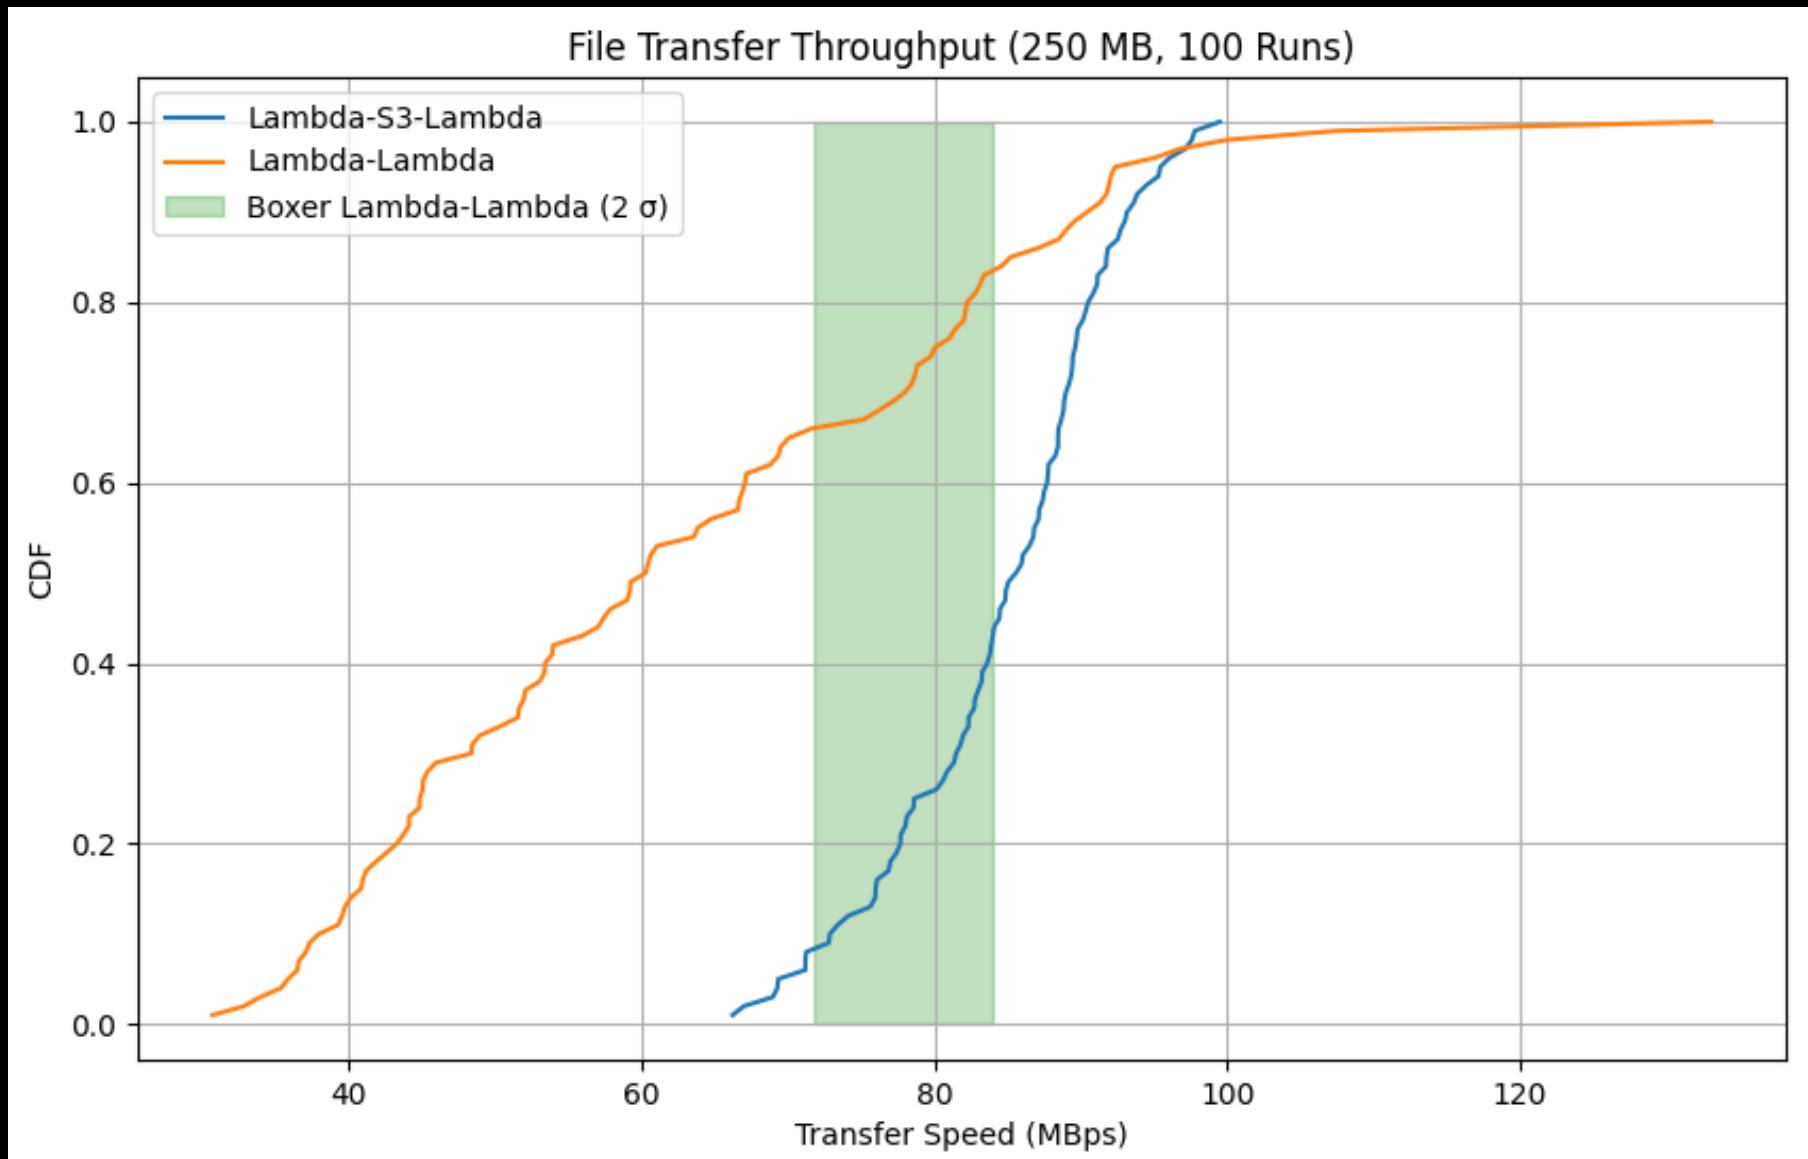
Reducers



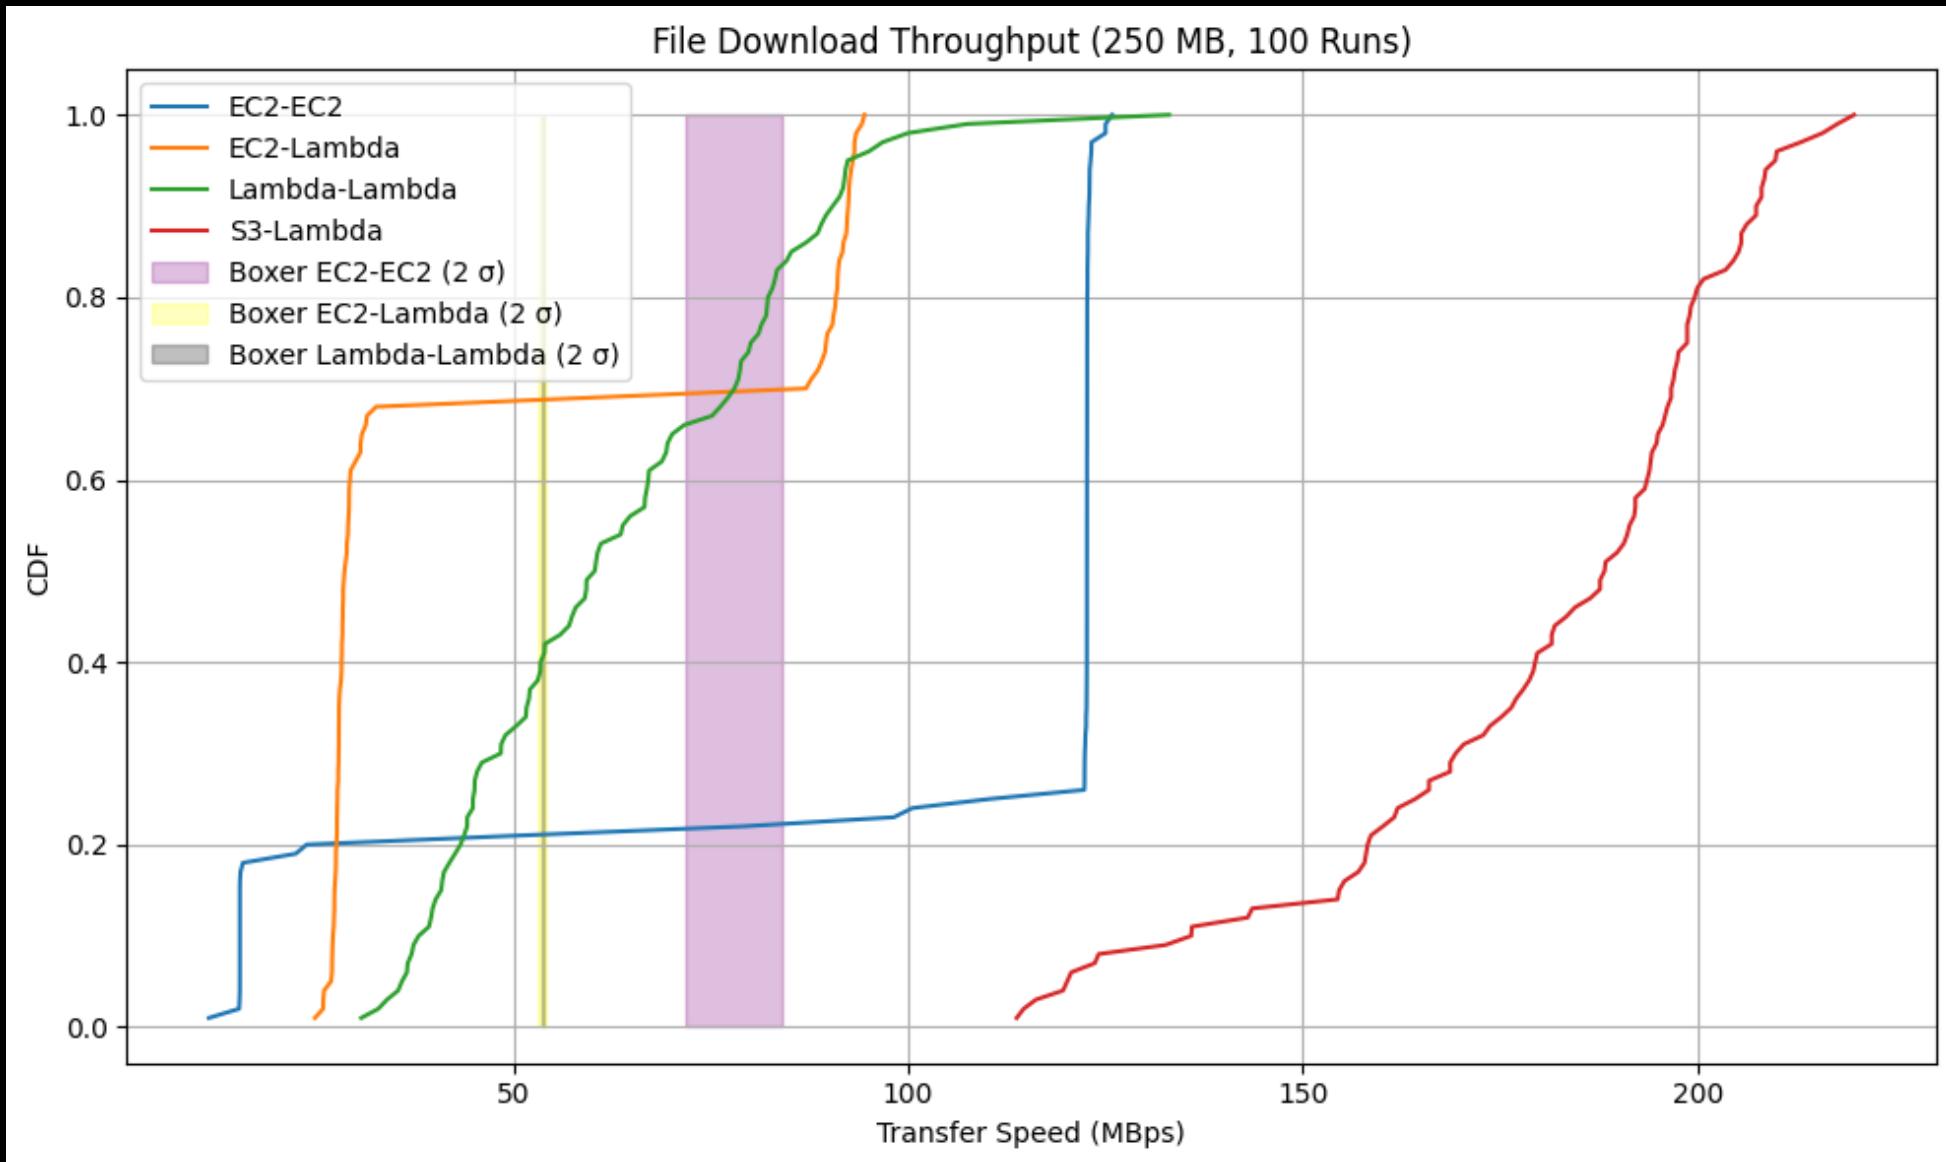
Performance Comparison



Performance Comparison

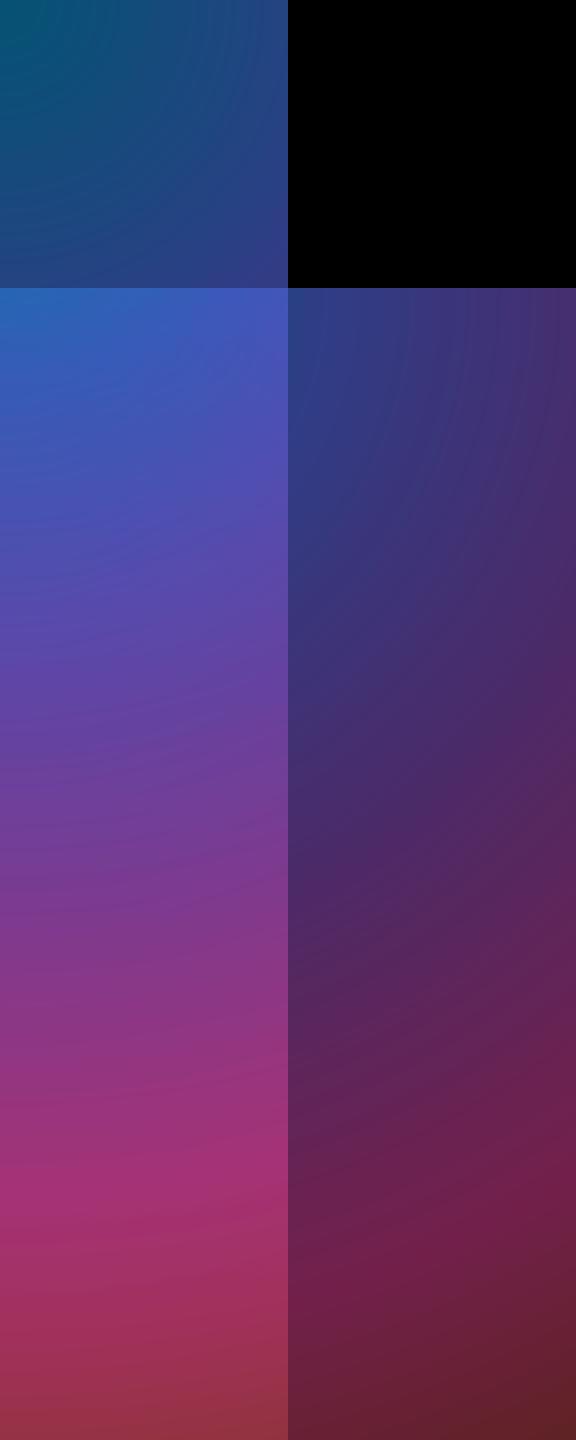


Performance Comparison



Battle Scars

- UDP hole punching is very sensitive to synchronization
 - All [INIT]s must occur within ~500ms of each other
 - Our lambdas **must** run in the same datacenter
- S3 performance improves with parallel uploads and downloads
 - However, rate limited to 5500 GET requests per second per file
- Reliable, in-order packet delivery over UDP is difficult
 - Don't try to implement your own protocol, use QUIC



Comparison with AWS EMR

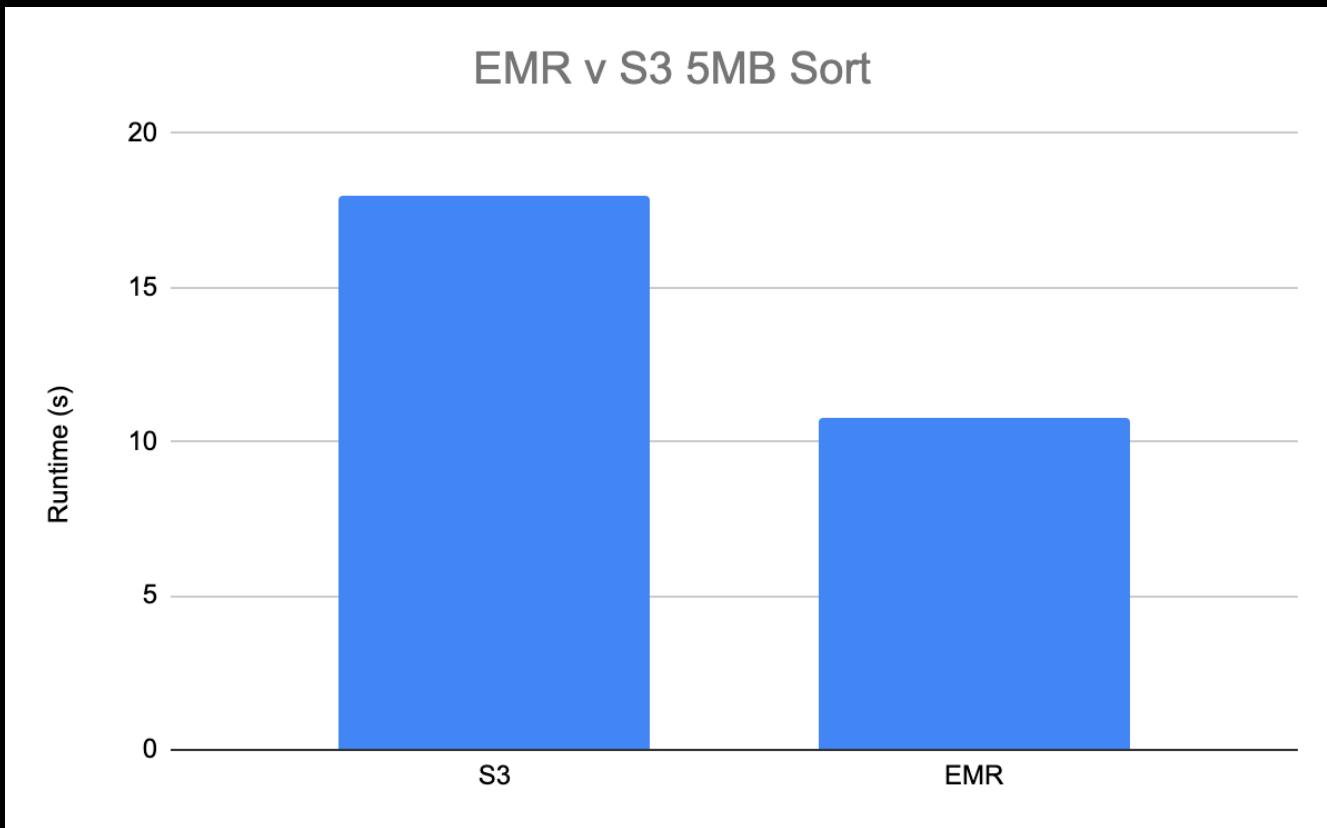
Amazon AWS EMR Serverless

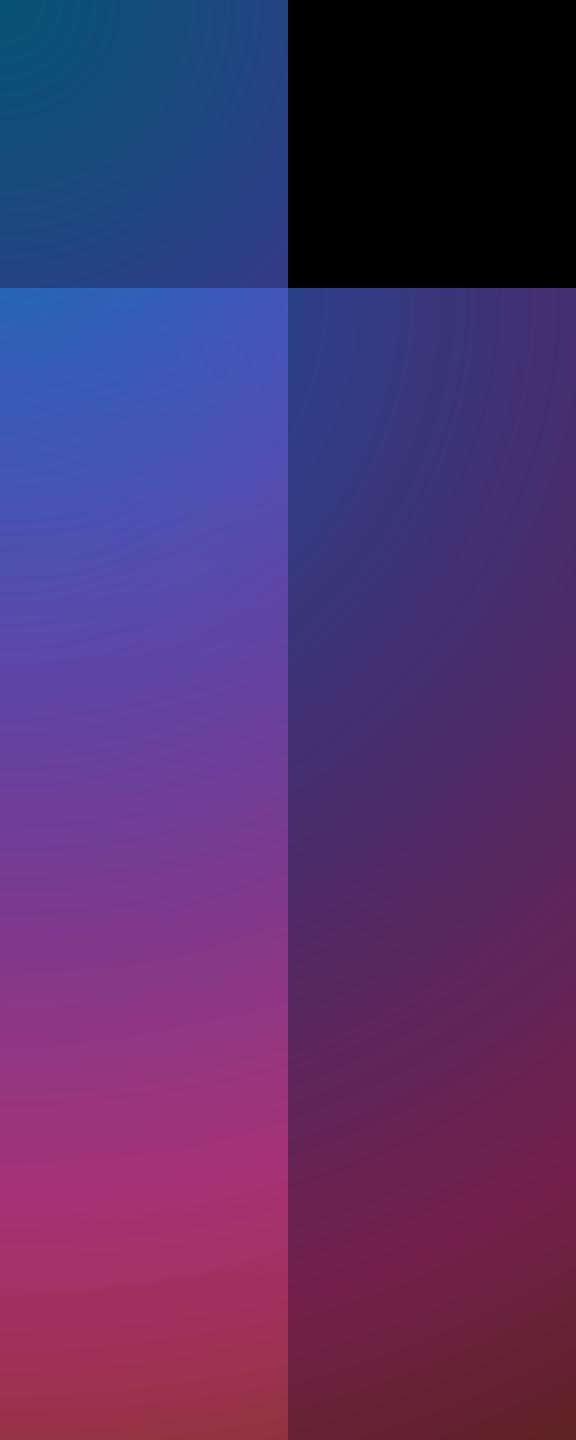
- Elastic MapReduce (EMR)
 - Managed cloud service
 - No need to provision
- Run jobs without provisioning
 - Determines required resources, allocates resources and releases them on completion

EMR Serverless Communication

- Communication
 - S3 used to communicate
 - Limited peer-to-peer communication under the hood, not configurable
 - Uses RPCs (remote procedure calls)
- Implements TCP/IP communication
 - Using the default `ElasticMapReduce-primary` role allows the core and task nodes to communicate with each other over TCP

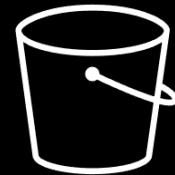
AWS EMR vs Lambda S3





Making it Resilient

Fault Tolerance Example

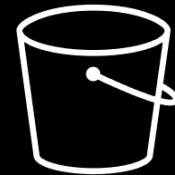


Mapper



Reducer

Fault Tolerance Example



1. Mapper
completes job

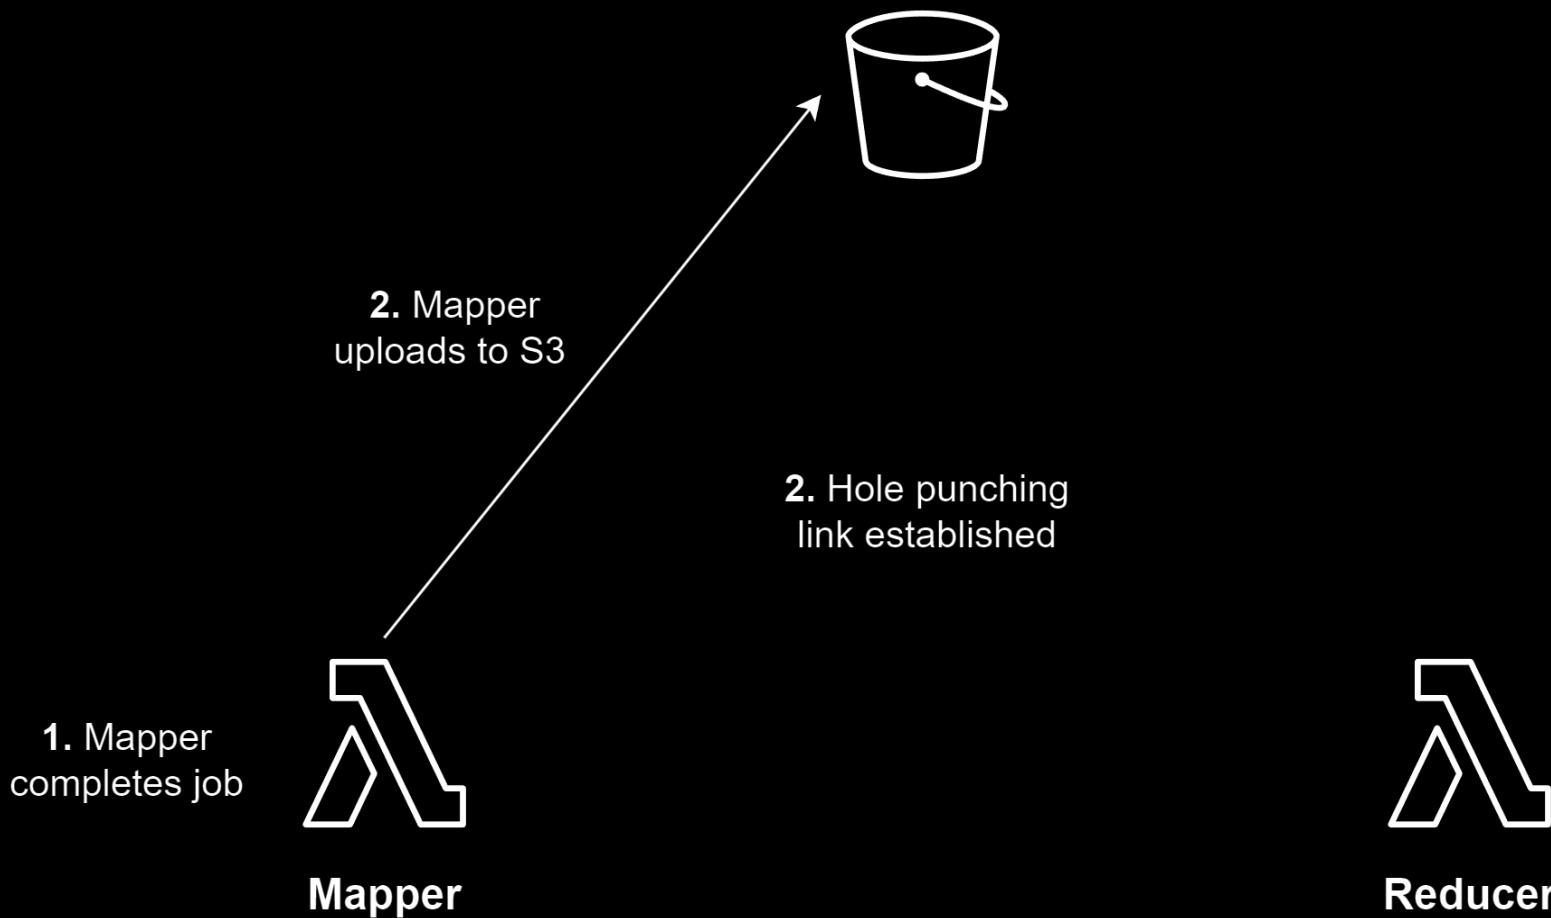


Mapper

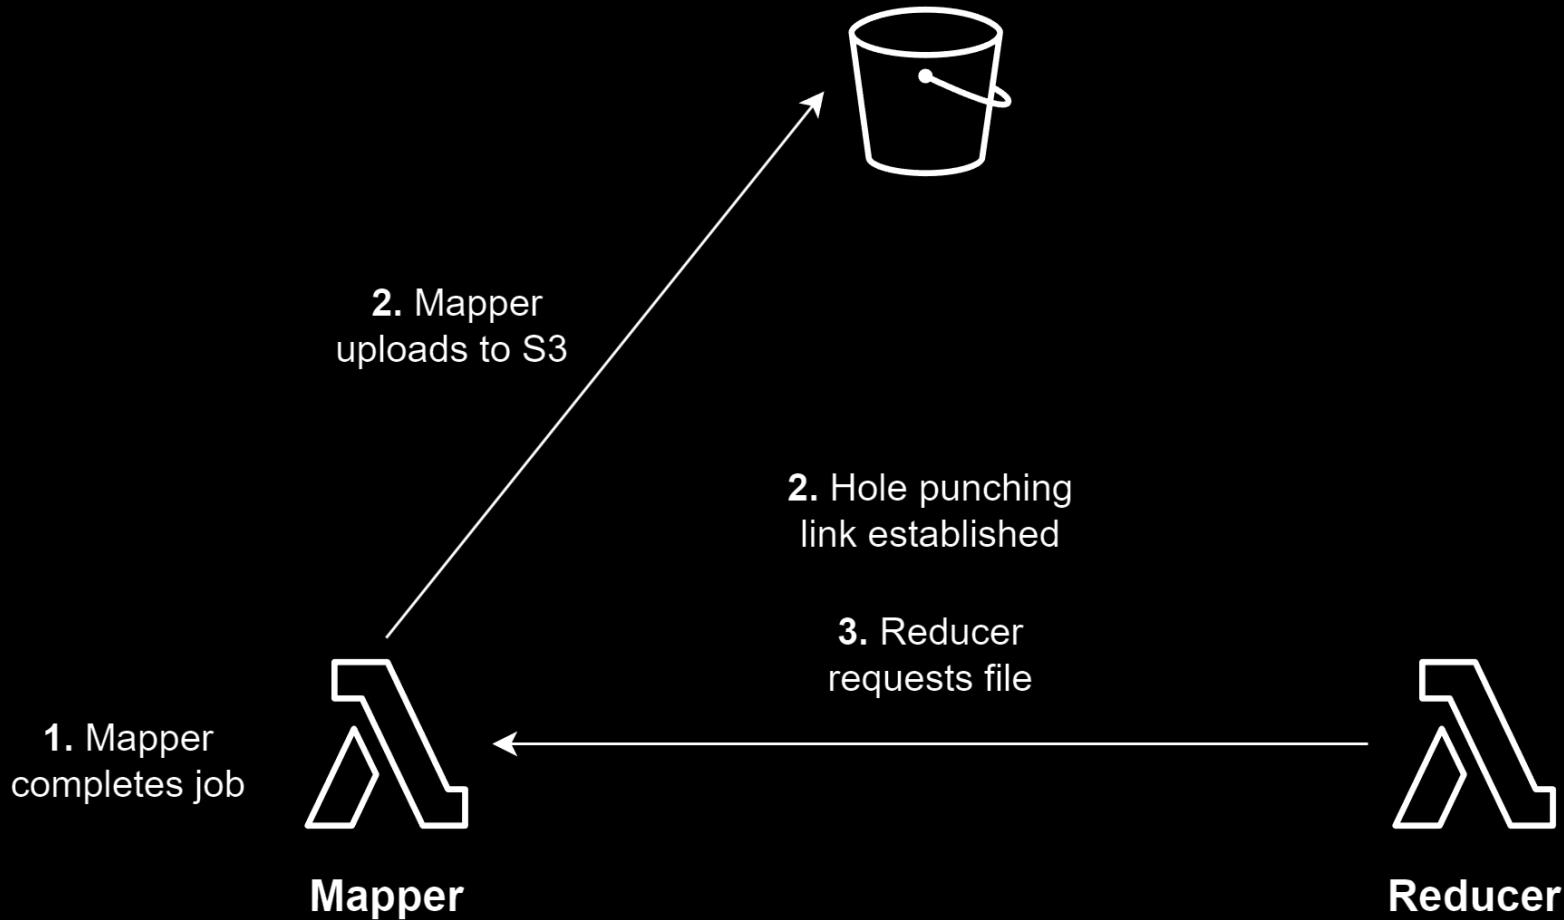


Reducer

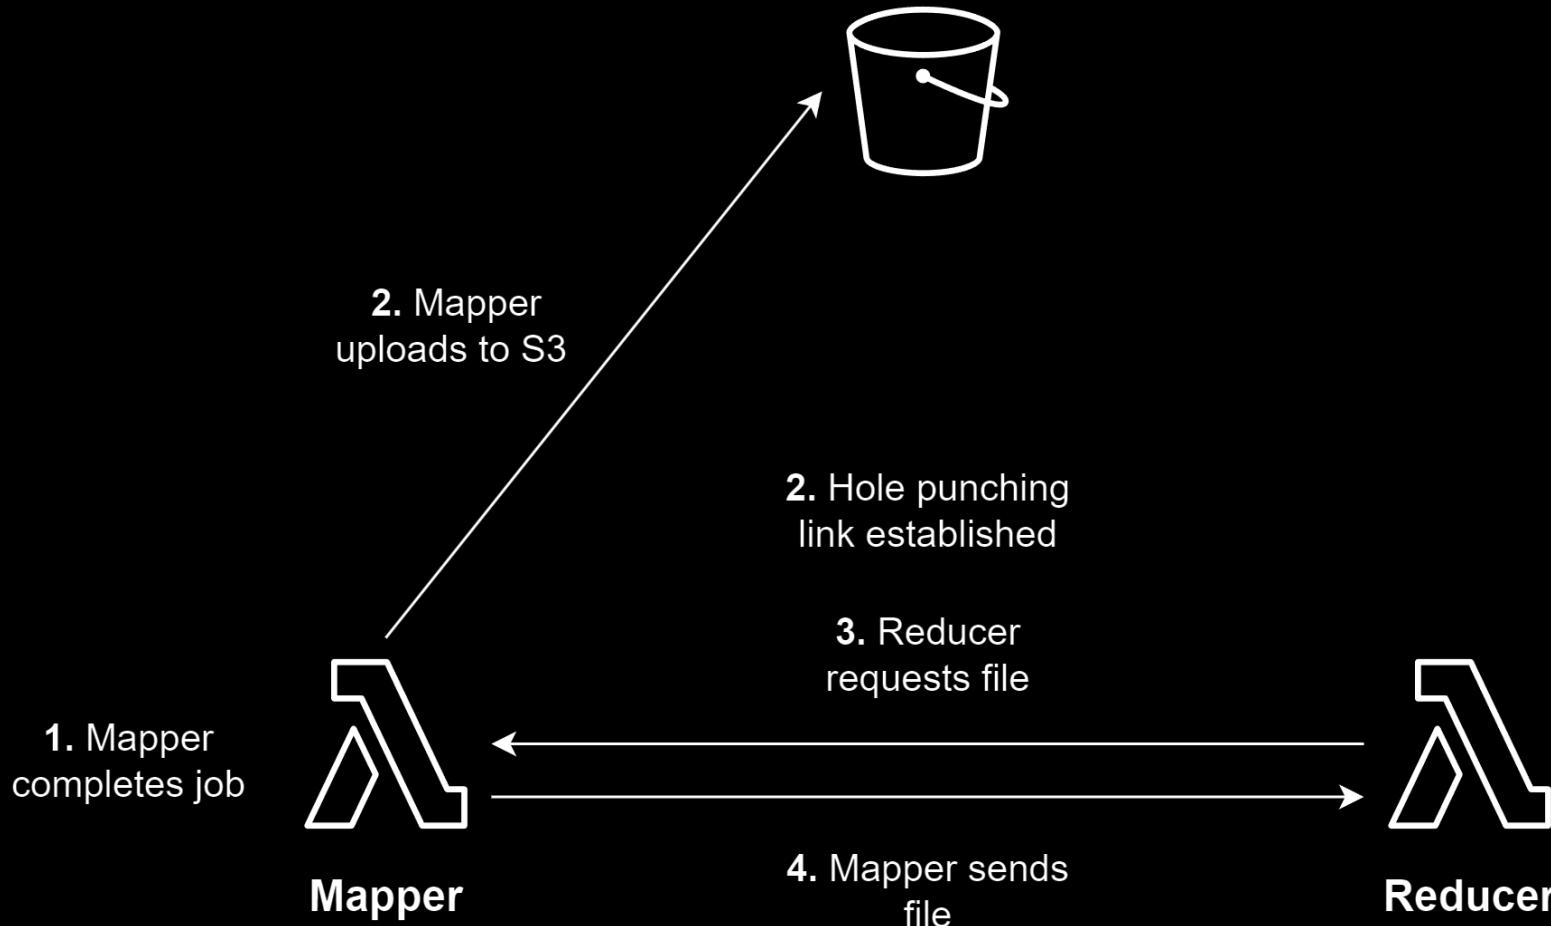
Fault Tolerance Example



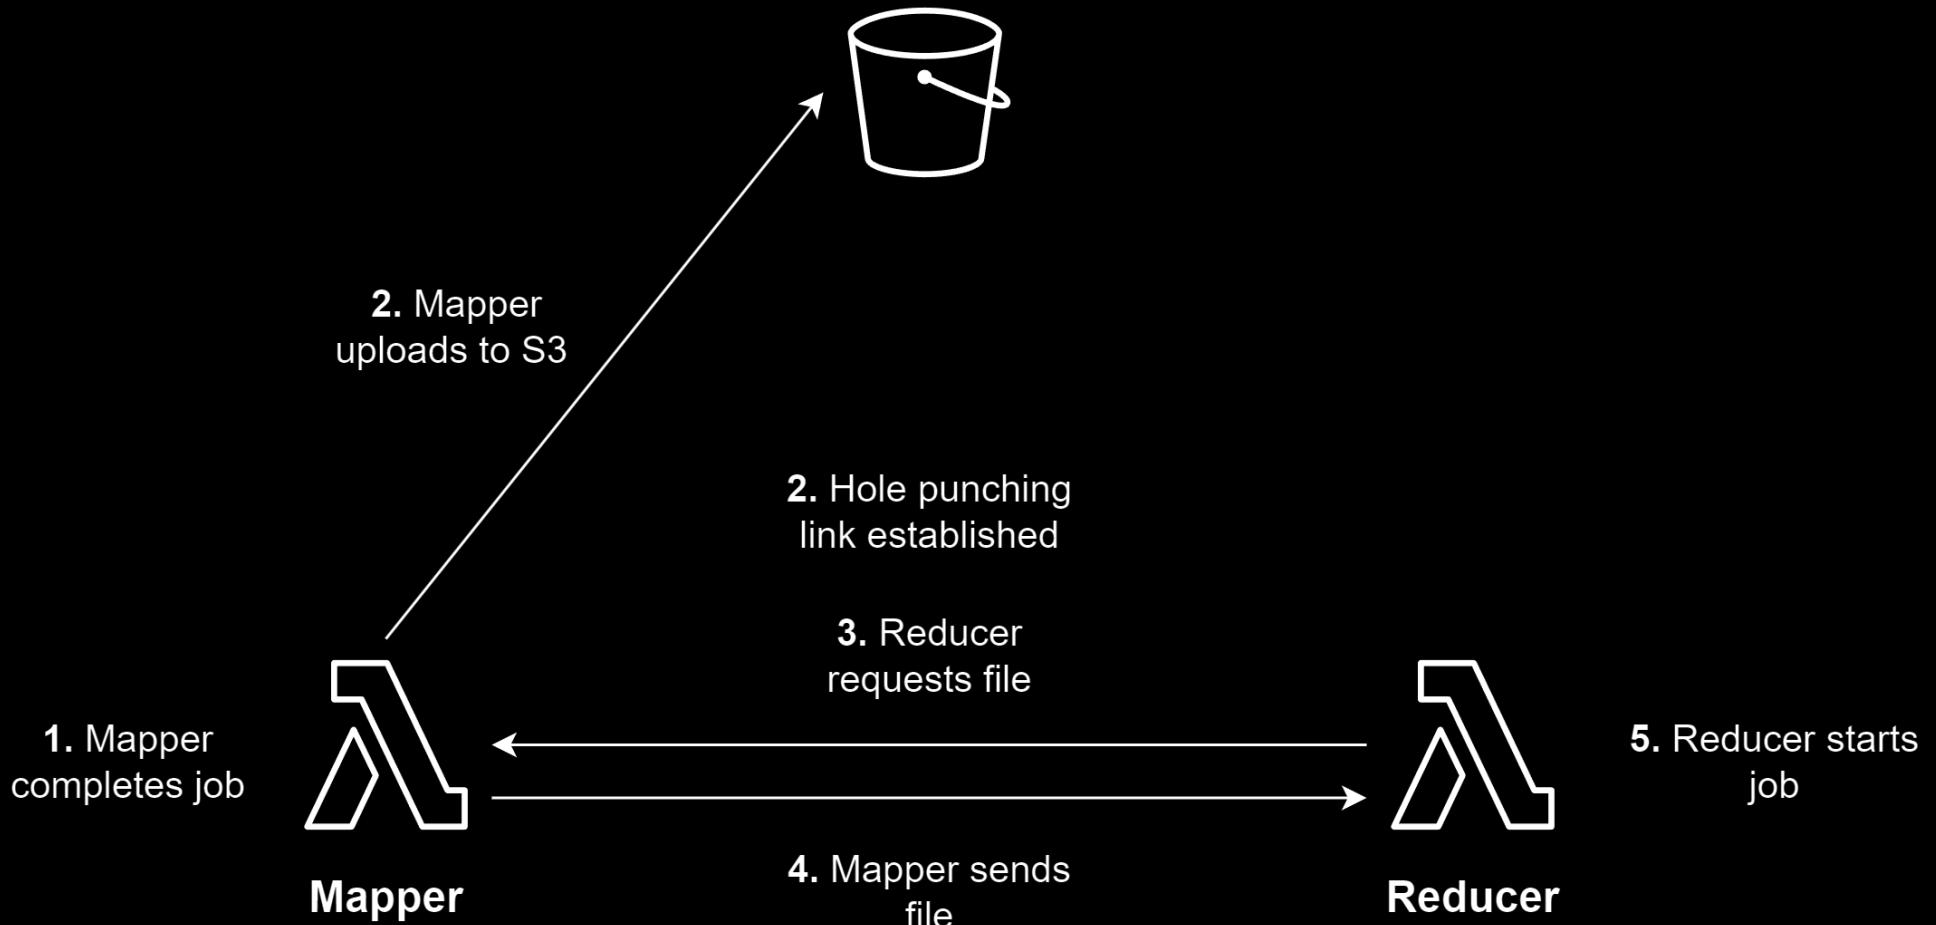
Fault Tolerance Example



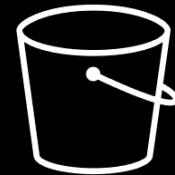
Fault Tolerance Example



Fault Tolerance Example



Fault Tolerance Example



1. Mapper
completes job

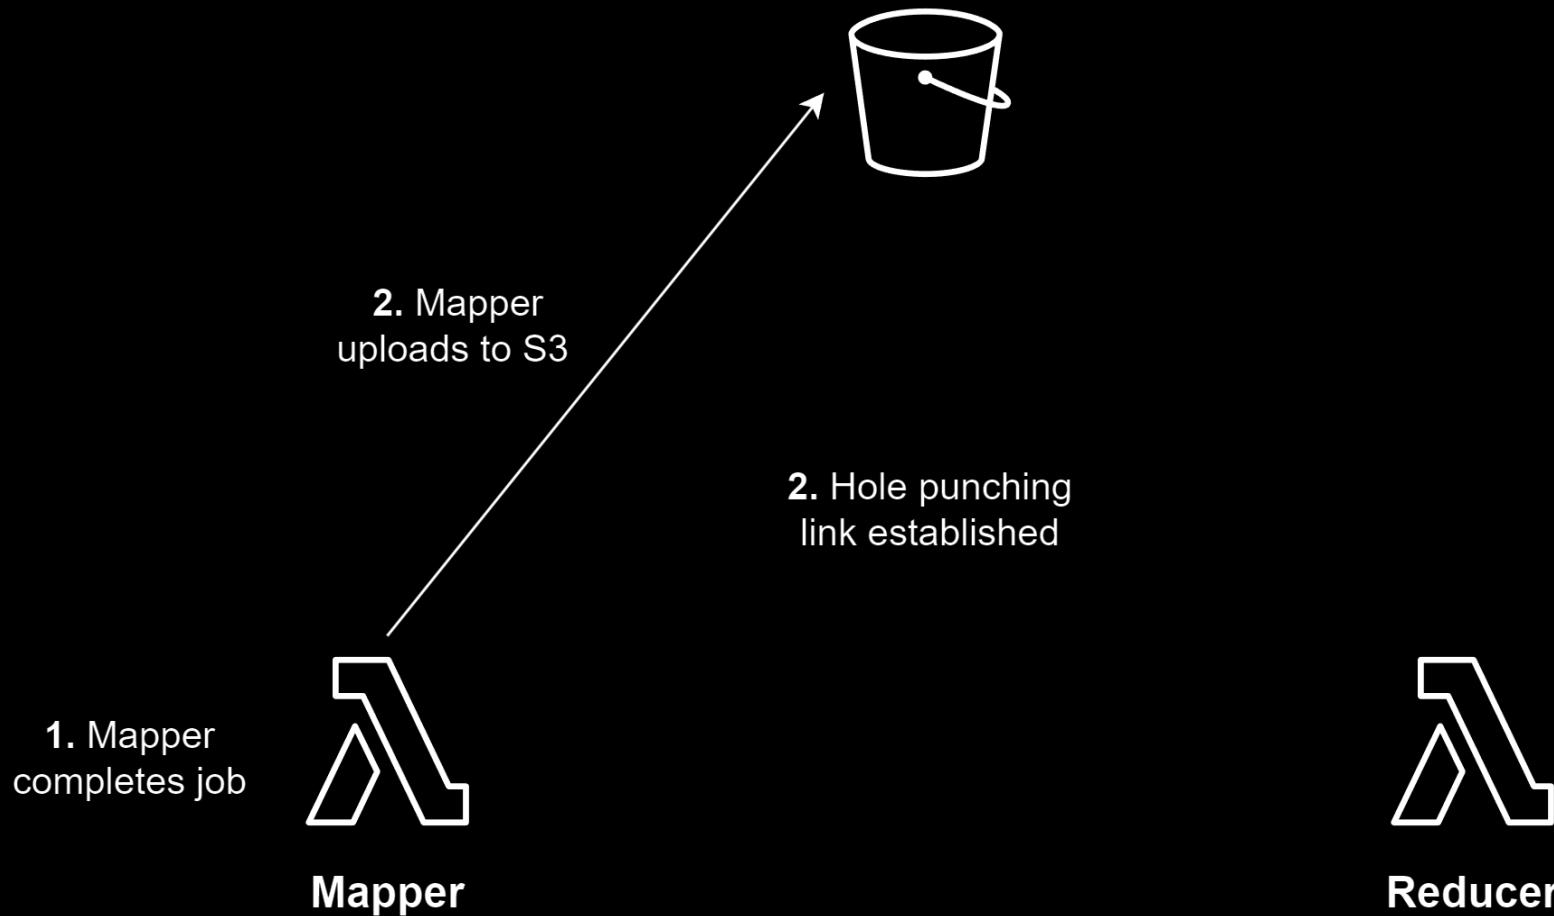


Mapper

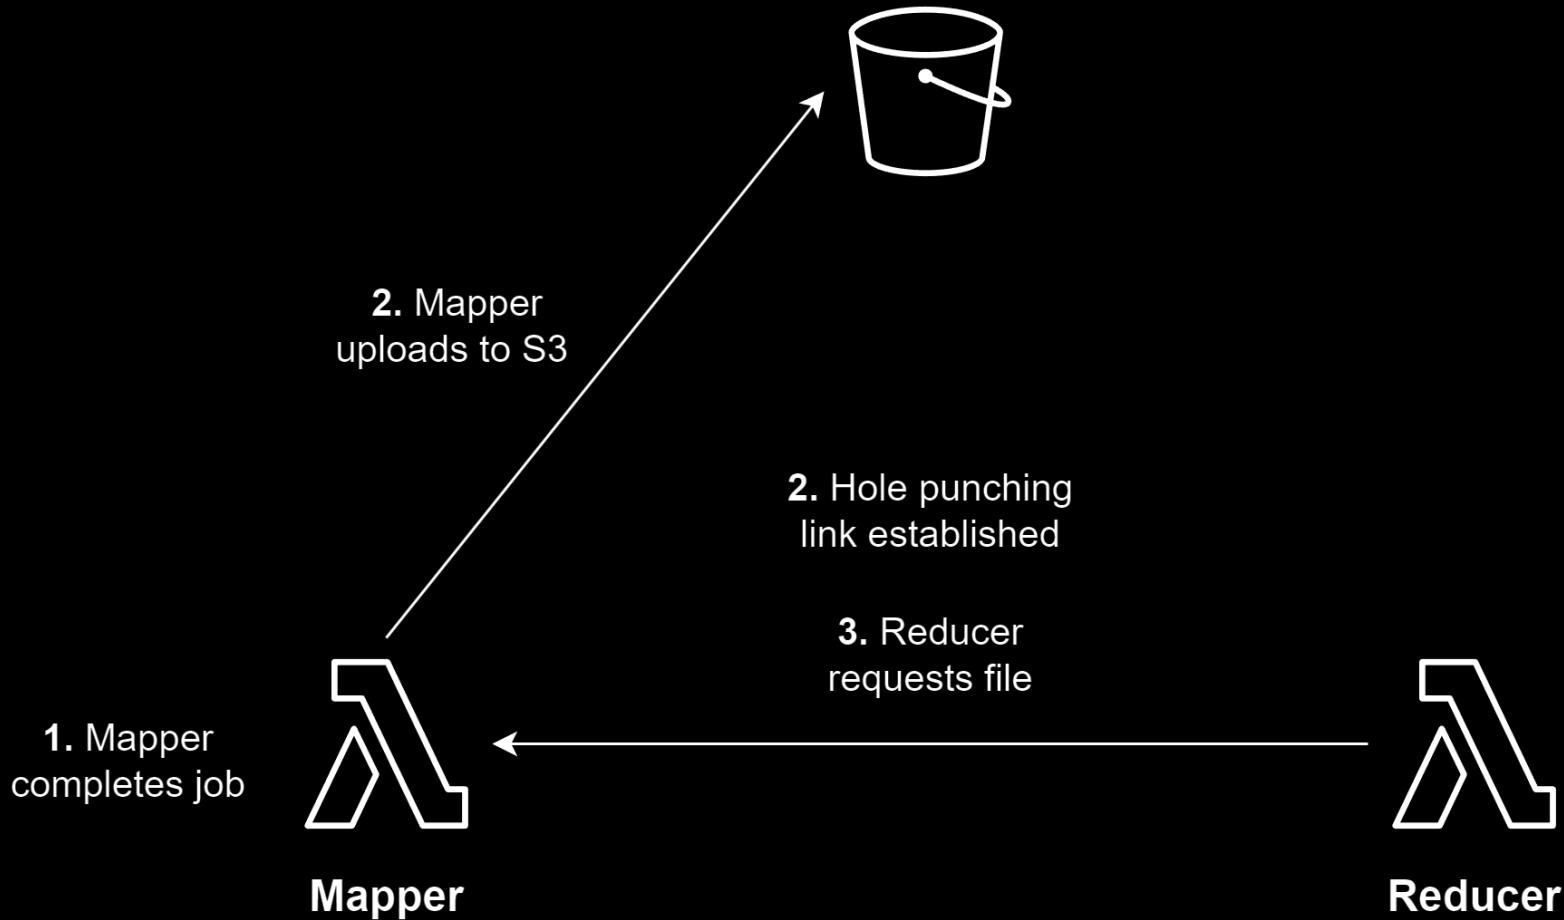


Reducer

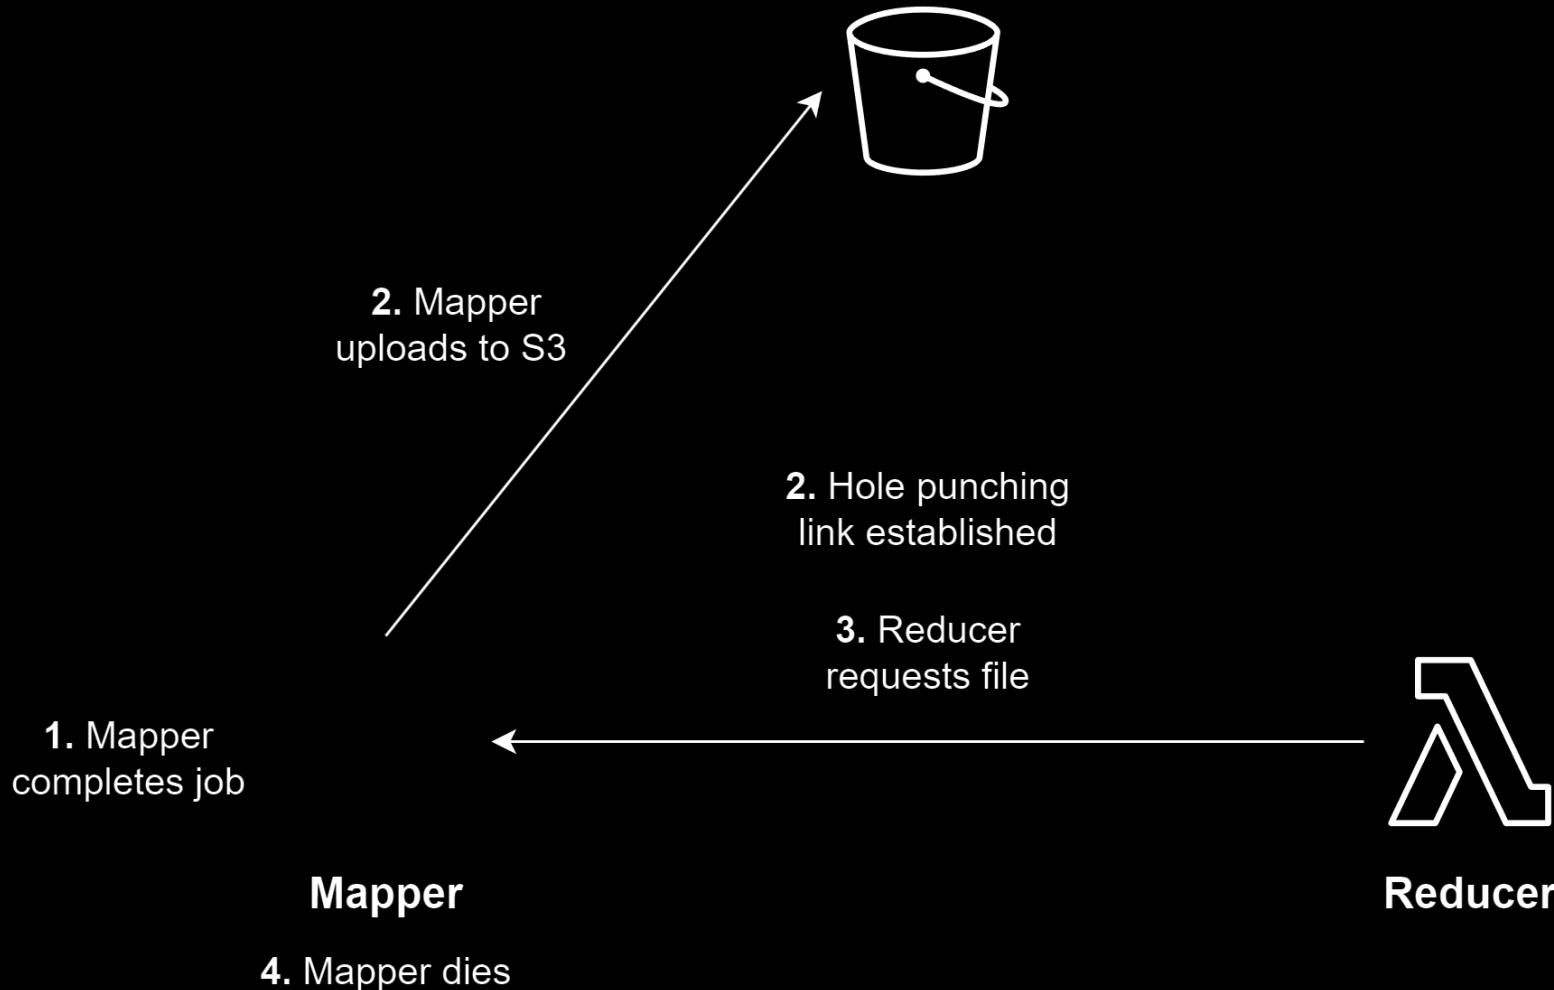
Fault Tolerance Example



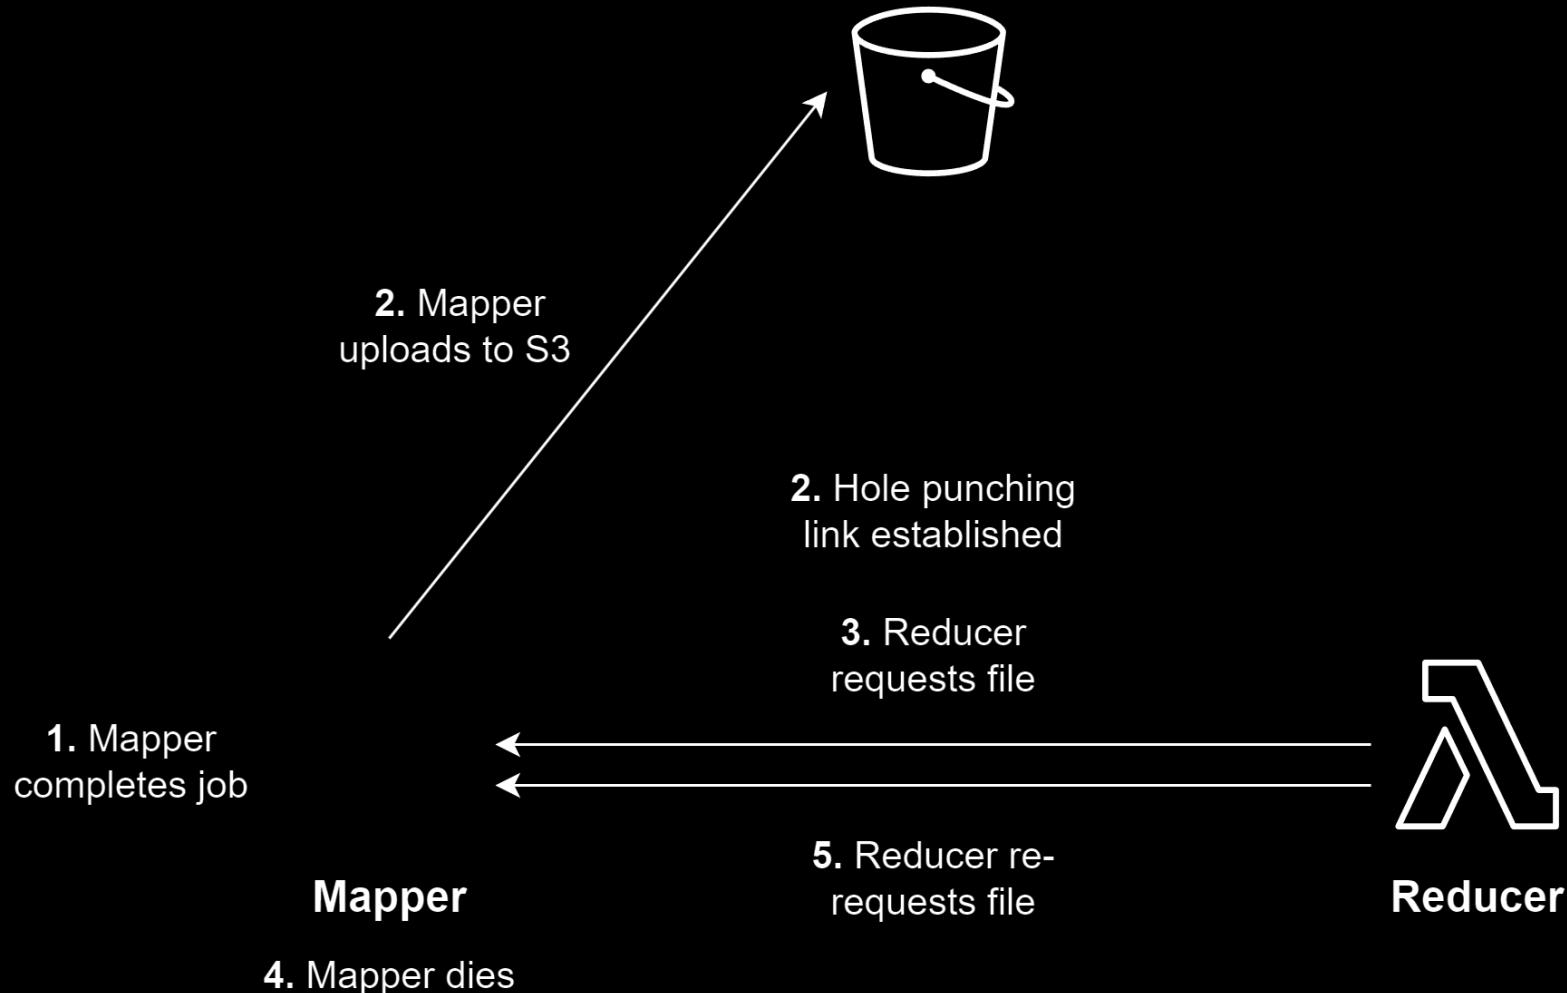
Fault Tolerance Example



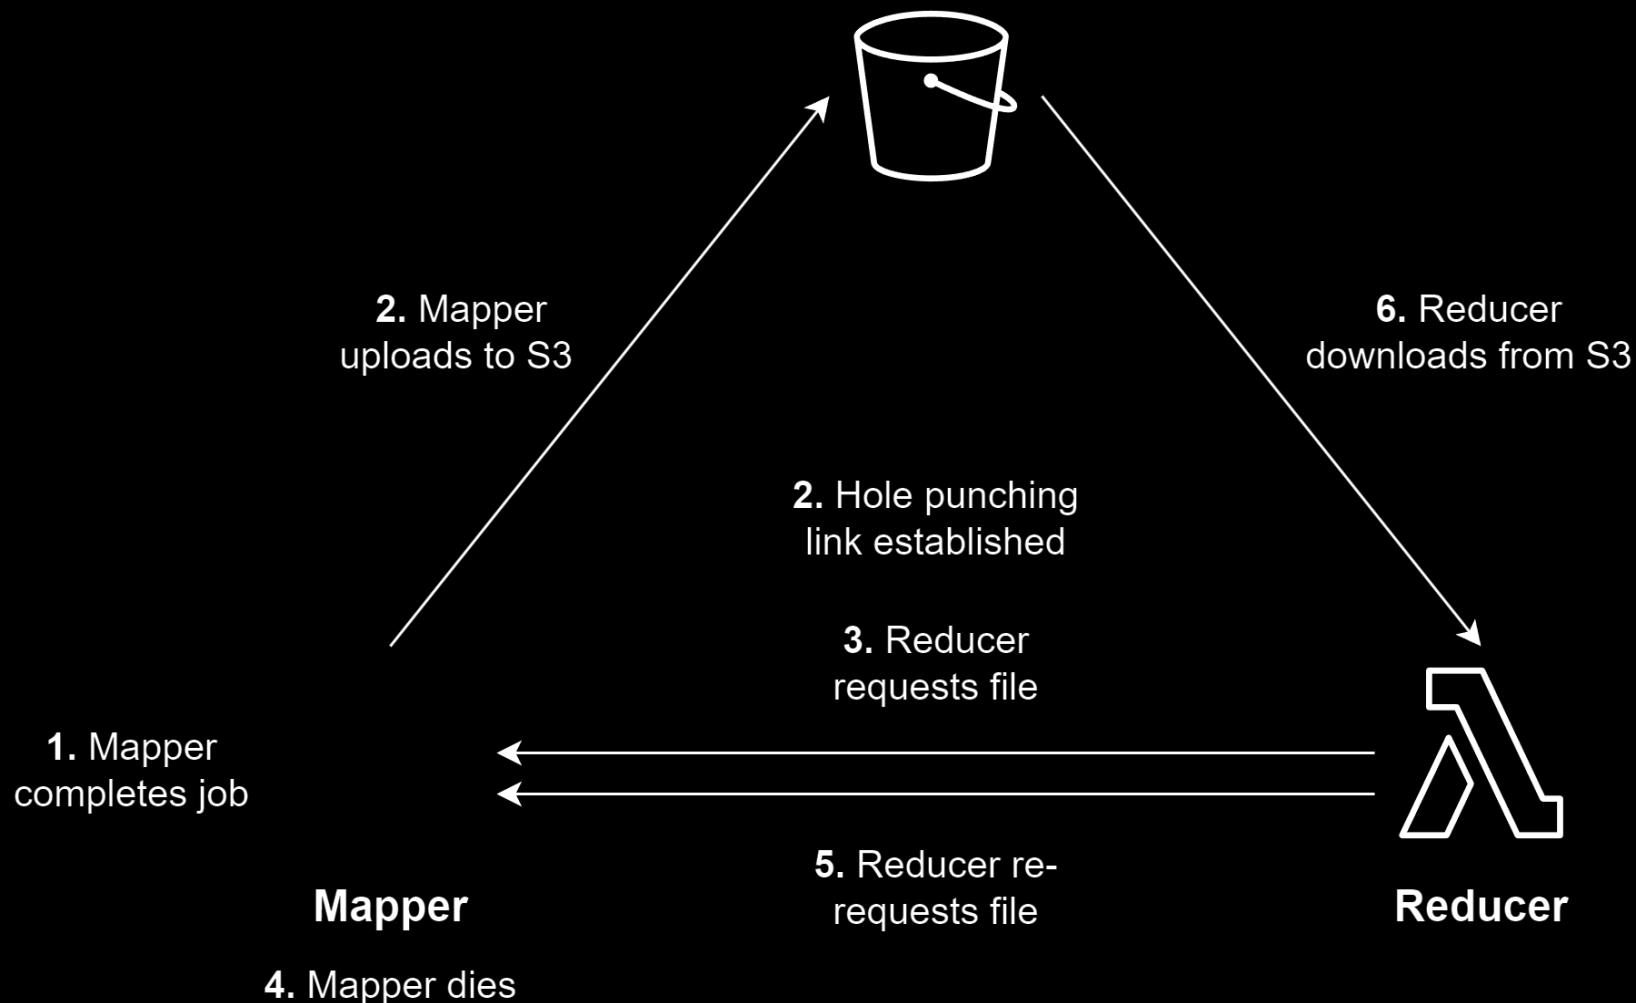
Fault Tolerance Example



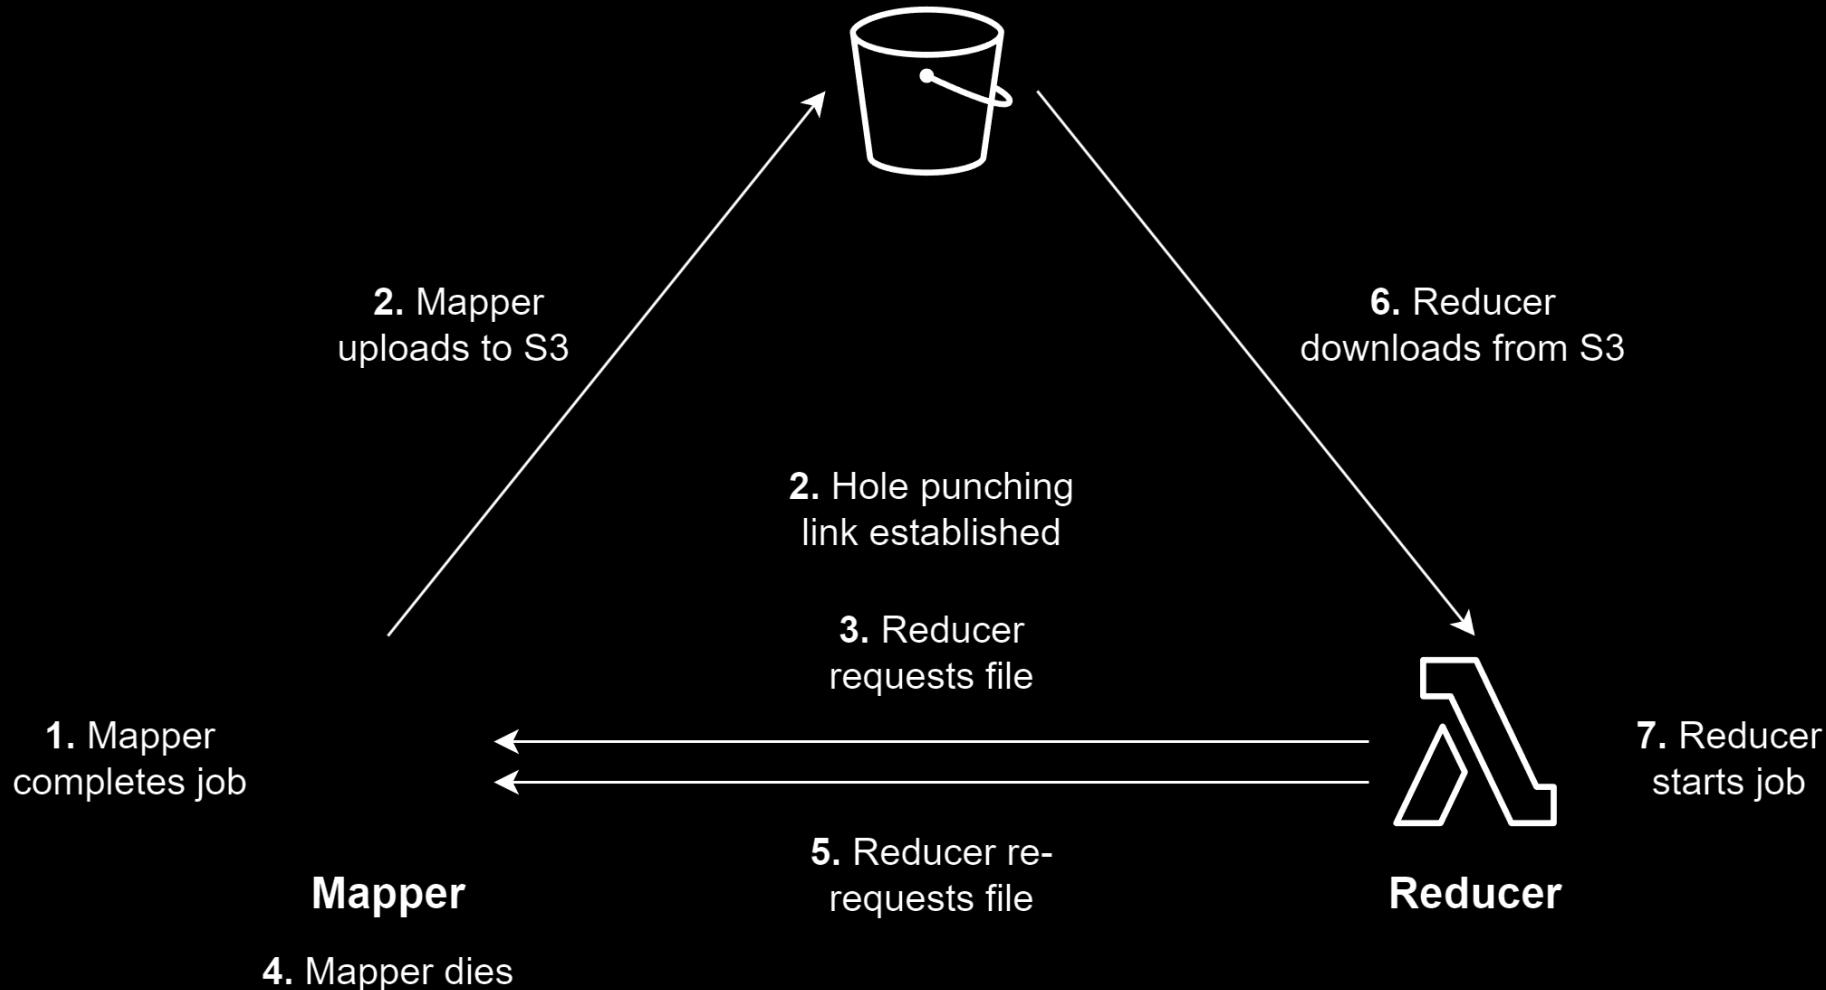
Fault Tolerance Example



Fault Tolerance Example



Fault Tolerance Example



Fault Tolerant Interfaces

Performance Centric

- Define a ***time to sync***
 - How long should the sender wait before writing to S3?
 - Navigates a durability-performance-cost tradeoff
- Define a ***time to live***
 - How long should the mapper stick around waiting for reducers to request data

Cost Centric

- Keep track of the **cost of compute** since the last checkpoint
- Sync once that value crosses a user defined **threshold**

Conclusion and Future Work

- Serverless is in need of **first-class network support**
- Through fault-tolerant communication, network support can be **aligned with the values of serverless**
- Providers like AWS **heavily optimize** their **paid services**
- We expect that the AWS paid tier would have better network performance
- Exploring different sized messages would be interesting
- Our implementation could be further optimized